Isometry Pursuit: Convex Recovery of Orthonormal Column Submatrices

Anonymous authors
Paper under double-blind review

Abstract

Isometry pursuit is a convex algorithm for identifying orthonormal column-submatrices of wide matrices.s It consists of a vector normalization followed by multitask basis pursuit. Applied to Jacobians of putative coordinate functions, it helps identity locally isometric embeddings from within interpretable dictionaries. We provide theoretical and experimental results justifying this method, including a proof with realistic assumptions that such isometric submatrices, should they exist, are contained within the obtained support. For problems involving coordinate selection and diversification, it offers a synergistic alternative to greedy and brute force search.

1 Introduction

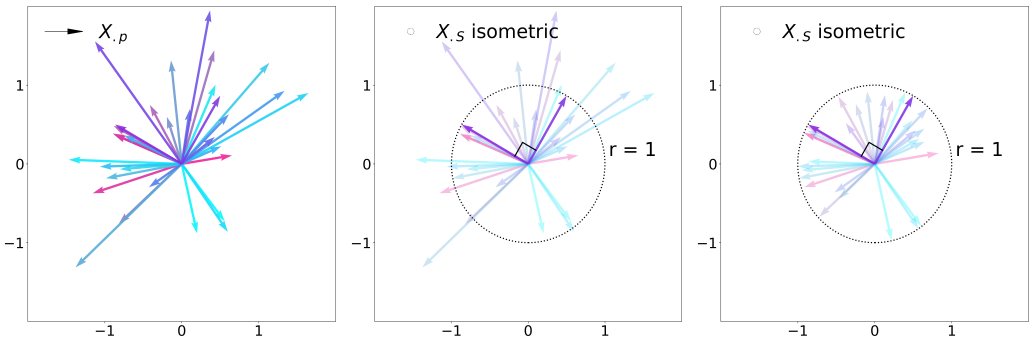
Many real-world problems may be abstracted as selecting a subset of the columns of a matrix representing stochastic observations or analytically exact data. This paper focuses on a simple such problem. Given a rank D matrix $X \in \mathbb{R}^{D\times P}$ with $P > D$, select a square submatrix $X_{.S}$ where subset $S \subset [P]$ satisfies $|S| = D$ that is as isometric as possible.

This problem arises in feature learning because while the coordinate functions of a given feature space may have no intrinsic meaning, it is often possible to generate a dictionary of interpretable features for consideration as potential parametrizing coordinates. When this is the case, selection of candidate interpretable features as coordinates can take the above form. Isometric embeddings - data mappings that compress high-dimensional distributions into low dimensions while preserving distances and angles on the data surface - have differentials that are isometric matrices. Thus, selecting isometric embeddings from without dictionaries of putative coordinates can be cast as a linear isometric submatrix coordinate selection problem. Identification of such coordinates generally aids mechanistic understanding, generative control, and statistical efficiency.

This paper takes an optimization perspective on the problem of selecting isometric submatrices. The background develops the the problem from the perspective of non-linear embeddings. However, the focus switches to a purely linear algebraic perspective in the remainder of the paper.

The main contribution of this paper is to show that an adapted version of the multitask basis pursuit feature learning algorithm in Koelle et al. [1] is a convex alternative isometry selection algorithm. This enables identification of low-dimensional coordinates within a Jacobian that preserve the original geometry of the data manifold under mapping given by the flow lines This algorithm - IsometryPursuit - is multitask basis pursuit applied to an appropriately normalized X selects orthonormal submatrices. Given vectors in $\mathbb{R}^D$, the normalization log-symmetrizes length and favors those closer to unit length, while basis pursuit favors those which are orthogonal. Our theoretical results formalize this intuition, and show the usefulness of IsometryPursuit as a trimming procedure prior to brute force search for diversification and interpretable coordinate selection. It is an alternative to greedy coordinate selection approaches such as those in Cai and Wang [2], Chen and Meila [3], Kohli et al. [4], Jones et al. [5].

[1]Work conducted outside of Amazon.



(a) Raw data (b) Isometric subset (c) Normalization step

Figure 1: 1a Raw data X 1b An isometric subset 1c Our method - vectors of length 1 maintain longest length after normalization

2 Background

Our goal is, given a matrix $X \in \mathbb{R}^{D \times P}$, to select a subset $S \subset [P]$ with $|S| = D$ such that $X_{.S}$ is as orthonormal as possible in a computationally efficient way. Figure 1a and 1b provide a visual depiction of this problem when $D = 2$ and $P = 30$. We motivate this problem from the perspective of non-linear embeddings.

2.1 Isometry

Isometries are a classical object in numerical linear algebra and geometry. They preserve distances and angles, and such can be used faithfully for subsequent human or algorithmic analysis. This motivates embedding algorithms such as Local Tangent Space Alignment [6], Multidimensional Scaling [7], and Isomap [8].

Definition 1 *A matrix $X \in \mathbb{R}^{D \times D}$ is an **isometry** or, equivalently, **unitary** if $XX^T = I_D$.*

Proposition 1 *The singular values $\sigma_1(X) \dots \sigma_D(X)$ are equal to 1 if and only if X is an isometry.*

Proposition 2 *If $P = D$, the component vectors $X_{.1} \dots X_{.D} \in \mathbb{R}^D$ form an isometry if and only if, for all $d, d' \in [D], \langle X_{.d}, X_{.d'} \rangle = \begin{cases} 1 & d = d' \\ 0 & d \neq d' \end{cases}$.*

Definition 2 *The **differential** of a smooth map $\phi : \mathcal{M} \to \mathcal{N}$ between D dimensional manifolds $\mathcal{M} \subseteq \mathbb{R}^B$ and $\mathcal{N} \subseteq \mathbb{R}^P$ is a map in tangent bases $x_1 \dots x_D$ of $T_\xi \mathcal{M}$ and $y_1 \dots y_D$ of $T_{\phi(\xi)} \mathcal{N}$ consisting of entries*

$$D\phi(\xi)_{ij} = \langle y_i, \frac{\partial \phi}{\partial x_j} \rangle (\xi) \tag{1}$$

Definition 3 *A map ϕ between D dimensional submanifolds is an **isometry at a point** $\xi \in \mathcal{M}$ if $D\phi(\xi)$ is an isometry.*

Definition 4 *A map ϕ between D dimensional submanifolds is **conformal at a point** $\xi \in \mathcal{M}$ if there exists a scalar $\lambda(\xi) > 0$ such that $D\phi(\xi)^T D\phi(\xi) = \lambda(\xi) I_D$.*

Proposition 3 *Let $\phi : \mathcal{M} \to \mathcal{N}$ be a map between D dimensional submanifolds.*

$$d_{\mathcal{N}}(\phi(\xi), \phi(\xi')) = d_{\mathcal{M}}(\xi, \xi'), \tag{2}$$

if and only if ϕ is smooth and is an isometry at every point [9, Problem 6-7].

Reference	Algorithm	Variable	Loss	Linearized	Domain	Identifies
Chen and Meila [3]	GREEDY	$S \in \binom{[P]}{M}$	$-\frac{1}{N}\sum_{n=1}^{N} \log \frac{\det\left(x_S^T x_S\right)^{1/2}}{\prod_{s\in S}\|x_S\|_2}(\xi_n)$	Yes	Dataset	Conformal
Kohli et al. [4]	GREEDY	$S \in \binom{[P]}{D}$	$\sup_{\xi',\xi''\in\mathcal{N}(\xi)\ \xi'\neq\xi''} r_S(\xi',\xi'') \sup_{\xi',\xi''\in\mathcal{N}(\xi)\ \xi'\neq\xi''} r_S(\xi',\xi'')^{-1}$ where $r_S(\xi',\xi'') := \frac{\|f_{.S}(\xi') - f_{.S}(\xi'')\|_2}{\|\xi' - \xi''\|_2}$.	No	Point	Conformal
Koelle et al. [1]	FISTA	$\beta \in \mathbb{R}^{N\times P\times D}$	$\frac{1}{2}\sum_{n=1}^{N}\|I_d - Df(\xi_n)\beta_{n\cdot\cdot}\|_F^2 + \lambda\sqrt{ND}\sum_{p=1}^{P}\left(\sum_{n=1}^{N}\sum_{d=1}^{D}\beta_{npd}^2\right)^{1/2}$	Yes	Dataset	Bijection

Table 1: Dictionary-based manifold learning losses. Note that we set the eigenfunction penalty to 0 for Chen and Meila [3] to emphasize the geometric condition

2.2 Coordinate selection

More recently, methods for selecting D coordinates from a dictionary $f_{1:P}$ of candidate features have emerged [4, 10**?**]. Loss functions for isometry vary along many axes. For example, functions may be measured at a single point [4] or over an entire dataset [10, 1, 3]. The loss function may make use of either linearized or non-linearized objectives. Objectives can be computed analytically or from numerical methods. Isometry may be relaxed as optimization target to the weaker notion of conformalness or embedding.

The dictionary-based coordinate search problem best subset selection problem is

$$\arg\min_{S\in\binom{[P]}{D}} \ell(f_{.S}) \tag{3}$$

where ℓ is some loss function that measures isometricness of a selected functional subset $f_{.S}$. Often, this problem is linearized via defining a loss at a point ξ as

$$\ell = \ell'(Df_{.S}(\xi)) \tag{4}$$

where ℓ' is now a loss function that measures the orthonormalness of the matrix $Df_{.S}(\xi)$. Non-linear objectives are often gradient estimators of linearized objectives, and so the distinction is smaller than it may appear.

Proposition 4 *The distortion loss function of f_S is equal to the condition number [11, 12] $\frac{\sigma_1(X)}{\sigma_D(X)}$ of $Df_{.S}(\xi)$ as $\mathcal{N}(\xi) \to \{\xi\}$.*

Proposition 5 *This is minimized by conformal maps.*

Furthermore, linearization can have computational advantages. For example, Koelle et al. [1] computes $Df_{.S}(\xi)$ via automatic differentiation.

2.3 Isometry losses

So far in this section, we have established the notion and desireableness of isometric embedding, and reviewed algorithms used for selecting column-submatrices that for conditions weaker than isometry. We now turn our review to the linear algebra literature.

Definition 5 *Sra [13]* The **S-divergence** $D(P, P')$ on positive definite matrices P, P' is given by

$$D(P, P') = (\log\det(\frac{P+P'}{2}) - \frac{1}{2}\log\det(PP'))^{\frac{1}{2}}. \tag{5}$$

Loss Function	Definition	Unique min.	At isometry	Unit length	Ortho-gonal
Distortion [4]	$\sigma_{\max}(X)/\sigma_{\min}(X) - 1$	No	Yes	No	Yes
Generic spectral loss	$\ell(X) = \sum_{d=1}^{D} g(\sigma_d(X))$	Depends on g	Depends on g	Depends on g	Depends on g
Normalized log-det volume	$-\sum_{d=1}^{D} \log \sigma_d \left(n_q^{\times P}(X)\right)$	No	No	Yes	No
Frob. pseudoinverse	$\|n_q^{\times P}(X)^\dagger\|_F - \sqrt{D}$	No	Yes	Yes	Yes
IsometryPursuit	$\|n_q^{\times P}(X)^\dagger\|_{1,2} - D$	Yes	Yes	Yes	Yes

Table 2: Matrix loss functions evaluated for isometry pursuit. We evaluate four key properties: whether the loss has a unique minimum, whether it is minimized specifically at orthonormal matrices (isometries), whether it prefers unit-length vectors over scaled versions, and whether it prefers orthogonal (orthonormal) structure.

Proposition 6 *If $X^T X$ is positive definite, then*

$$D(X^T X, I_D)^2 = \sum_{d=1}^{D} \log \frac{\sigma_d(X) + \sigma_d(X)^{-1}}{2}. \tag{6}$$

Proposition 7 *If $X^T X$ is positive definite, then $D(X^T X, I_D)^2$ is uniquely minimized by isometric matrices.*

See [14] for a more extensive discussion of other matrix divergences. In Section 3, we will examine matrix loss functions capable of selecting isometric submatrices based on their singular values from a generic perspective.

2.4 Sparse convex coordinate selection

Equipped with a function ℓ that is minimized by an isometry, objective functions like 3 may be optimized in a variety of ways. Brute force algorithms consider all possible solutions are guaranteed to find the global minimum, but have the often prohibitive time complexity $O(C_\ell P^D)$ where C_ℓ is the cost of evaluating ℓ. In contrast, greedy algorithms consist of iteratively adding one element at a time to S. This algorithms have time complexity $O(C_\ell PD)$ and so are computationally more efficient than brute force algorithms, but can get stuck in local minima. Formal definitions are given in Section 6.1.

Sometimes, it is possible to introduce an objective which convexifies problems of the above form [15, 16, 17].

$$\arg \min_{\beta \in \mathbb{R}^P : y = x\beta} \ell(\beta) \tag{7}$$

to the overcomplete regression problem $y = x\beta$ are a classic example [18]. Overcomplete regression problems are distinguished by the fact that $P > D$ and so the solution β is not unique. When $f(\beta) = \|\beta\|_0$, this non-convex problem can be solved by any subset of sufficient rank. However, when $f(\beta) = \|\beta\|_1$, the problem is not only more geometrically restrictive, but it is also only convex and may be solved efficiently e.g. via interior-point methods. See Schneider and Tardivel [20] for a recent discussion.

3 Method

The contributions of this paper is that we introduce an ℓ uniquely minimized by isometries that is convex in its parameters, induces sparsity, and is more appropriate than S-divergence when an isometric submatrix is not present in that it always favors vectors of unit length. This ℓ is a multitask inspired by the group basis pursuit [10, 21] which we call differently since grouping is dependent only on the structure of matrix-valued response variable I_D. We show that this approach, used to select independent dictionary elements in Koelle et al. [10, 1], is in fact capable of usefully solving the more specific problem of selecting pointwise isometries from a dictionary.

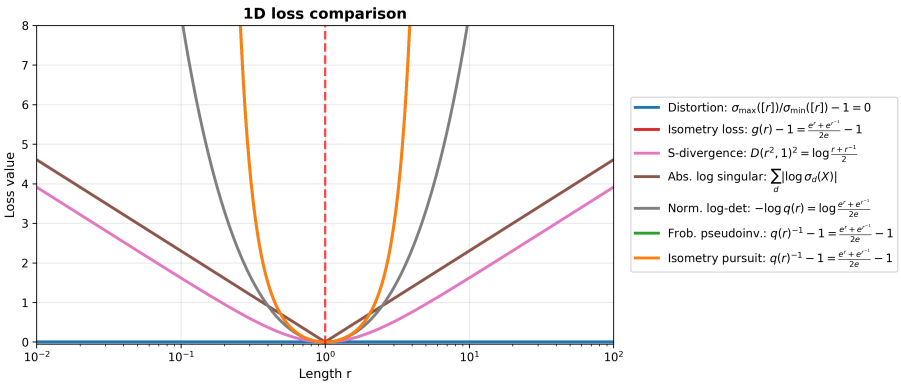


Figure 2: One-dimensional candidate loss curves for $X = [[r]]$, showing loss behavior versus vector length. The red vertical dashed line marks $r = 1$.

3.1 The problem with singular values

(8)

(9)

$\ell_S(X) = D(X^T X, I_D)$ is minimized uniquely by orthonormal matrices and invariant under rotation. However, it is plainly not the only function of X satisfying these properties. Consider that function

$$\ell : \mathbb{R}^{D \times P} \to \mathbb{R}^+ \tag{10}$$

$$X \mapsto \sum_{d=1}^{D} g(\sigma_d(X)) \tag{11}$$

where

$$g : \mathbb{R}^+ \to \mathbb{R}^+ \tag{12}$$

$$t \mapsto \log \frac{t + t^{-1}}{2}. \tag{13}$$

Clearly, any g uniquely minimized at 1 will provide a suitable rescaling of singular values. However, there are some nice properties we may also want. Monotonicity away from 1. $g(t) = t(t^{-1})$. See for example Cichocki et al. [14] for a variety of log-polynomial singular value divergences on positive definite matrices satisfying these property. We provide a heuristic answer to the first question and a more rigorous to the second. Intuitively, the problem with singular values seems to be that they can blow up when X is rank deficient and so aren't very good for partially correct answers. On the other hand, the log-polynomial loss grows quite slowly with error and so is overly forgiving.

The deep issue is more subtle. Essentially, we want a loss that grows the further from isometry we are. Among the conditions for satisfying this are that given a set of angles between vectors we want the loss to be less when the vectors are length 1. This is an issue for singular-value based losses, since two colinear vectors of length 1 will increase a singular value to $\sqrt{2}$. In this way, as one of the vectors is shortened, the singular value passes through 1. What of the other singular values? It was quite small, and it now has gotten a little bit smaller, but not commesurately with the decrease in loss from the top singular value.

Put mathematically, we want

$$\frac{\partial \ell}{\partial \|x_{.d}\|_2} \begin{cases} > 0 & : \ \|x_{.d}\|_2 > 1 \\ = 0 & : \ \|x_{.d}\|_2 = 1 \\ < 0 & : \ \|x_{.d}\|_2 < 1 \end{cases} \tag{14}$$

for all X. What this property says is that, as a column-set becomes more normal, the loss is lower. This property is plainly of interest for isometry selection.

What is also of interest is the following log-symmetry property.

Unfortunately, we show the following proposition.

Proposition 8 *There does not exist a g such that $\ell = \sum_{d=1}^{D} g(\sigma_d(X))$ such that ℓ satisfies Property and*

Property 1

$$g(t) = g(1/t) \tag{15}$$

A suitable counterexample is shown in Section 6.2.2. The issue basically is that colinear vectors punish singular value based losses by pushing larger singular values away from 1. The issue can also be seen visually in Figure **??**, where the level sets of the log det losses are not tangent to length at $r = 1$, meaning that given an angle loss can actually be decreased by shrinking

Although starting from singular values is bad in the sense of Proposition , other losses do satisfy this property. These losses arise from the intuition that we can shrink vectors away from length 1 in a manner similar to how we selected g. Then, we can apply some loss function to the rescaled matrix that satisfies Property 8.

In fact, two fundamentally different loss functions first come under consideration. The determinant of the rescaled matrix will be maximized at orthonormality. Determinant is monotonic with product of singular values. So now basically we can do something with product of the singular values rather than their log sum.

Why is product bad?

Sum of singular values is the trace. So will the sum of its squared singular values. On the other hand we can apply a method-of-frames penalty. Also note that trace of A transpose A is in fact the nuclear norm.

These divergences raise the question of what is the most reasonable form of g, and whether starting with singular values is a good way to define an isometry loss.

$$g : \mathbb{R}^+ \to \mathbb{R}^+ \tag{16}$$

$$t \mapsto \frac{e^t + e^{t^{-1}}}{2e}. \tag{17}$$

By Proposition 1, we can see that ℓ is uniquely maximized by orthonormal matrices. Moreover, g is convex, and $\ell(X^{-1}) = \ell(X)$ when X is invertible. Figure 2 gives graphs of candidate losses when $D = 1$, and Figure 3 gives the analogous landscapes for $D = 2$. Figure 4 gives the analogous two-parameter view for $D = 3$ equicorrelated matrices.

Our ground truth objective is therefore the best subset selection problem

$$\arg \min_{S \in \binom{[P]}{d}} \ell(X_{\cdot S}). \tag{18}$$

Regardless of the convexity of ℓ, brute combinatorial search over $[P]$ is inherently non-convex. This motivates our alternative formulation.

3.2 Normalization

Since basis pursuit methods tend to select longer vectors, selection of orthonormal submatrices requires normalization such that both long and short candidate basis vectors are penalized in the subsequent regression. We introduce the following definition.

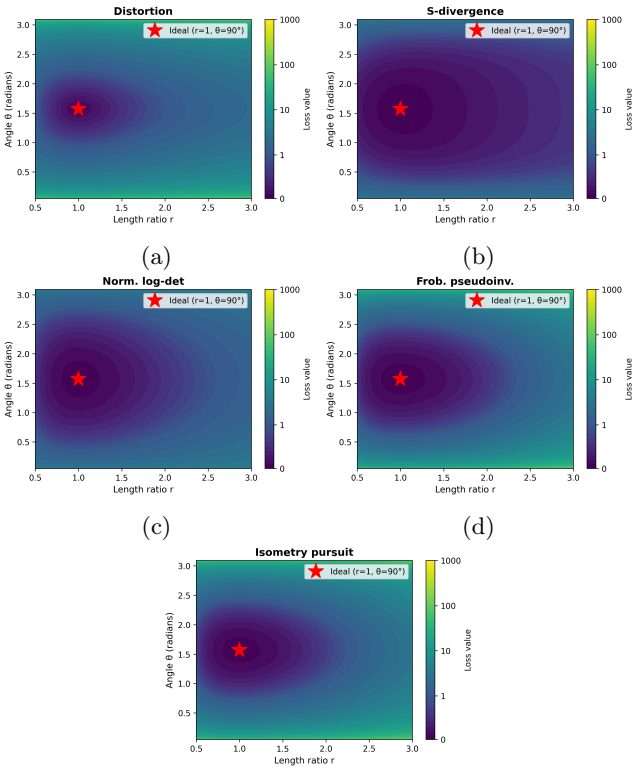


Figure 3: Two-dimensional loss function comparison for isometry pursuit. The 2D loss landscapes are for 2×2 matrices parametrized by length ratio r and angle θ with shared colorbars. S-divergence is shown as the representative generic spectral loss. The red star in each 2D plot marks $(r = 1, \theta = \pi/2)$ (ideal). The panels visualize whether each candidate loss prefers unit-length and orthogonal structure at the ideal point. See Table 2 for mathematical definitions.

Definition 6 (Symmetric normalization) *A function* $q : \mathbb{R}^D \to \mathbb{R}^+$ *is a symmetric normalization if*

$$\arg\max_{v \in \mathbb{R}^D} \; q(v) = \{v : \|v\|_2 = 1\} \tag{19}$$

$$q(v) = q\left(\frac{v}{\|v\|_2^2}\right) \tag{20}$$

$$q(v_1) = q(v_2) \; \forall \; v_1, v_2 \in \mathbb{R}^D : \|v_1\|_2 = \|v_2\|_2. \tag{21}$$

We use such functions to normalize column-vector length in such a way that vectors of length 1 prior to normalization have longest length after normalization and vectors are shrunk proportionately to their deviation from 1. That is, we normalize vectors by

$$n : \mathbb{R}^D \to \mathbb{R}^D \tag{22}$$

$$v \mapsto q(v)\frac{v}{\|v\|_2} \tag{23}$$

and matrices by

$$w : \mathbb{R}^{D \times P} \to \mathbb{R}^{D \times P} \tag{24}$$

$$X_{\cdot p} \mapsto n(X_{\cdot p}) \; \forall \; p \in [P]. \tag{25}$$

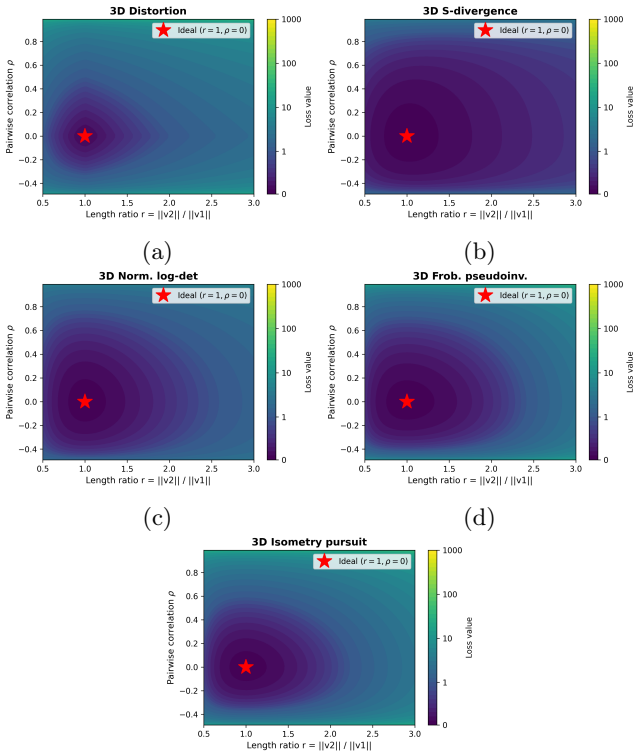


Figure 4: Loss function comparison for 3D equicorrelated matrices, shown as two-dimensional landscapes. Each matrix consists of three vectors in $\mathbb{R}^3$ with lengths $(1, r, 1)$ and common pairwise correlation ρ. The horizontal axis varies the length ratio $r = \|v_2\|_2/\|v_1\|_2$, while the vertical axis varies ρ. S-divergence is shown as the representative generic spectral loss. The red star marks the ideal orthonormal setting $(r, \rho) = (1, 0)$.

We choose q as follows.

$$q : \mathbb{R}^D \to \mathbb{R}^+ \tag{26}$$

$$v \mapsto \frac{2e}{e^{\|v\|_2} + e^{\|v\|_2^{-1}}}. \tag{27}$$

Besides satisfying the conditions in Definition 6, this normalization has some additional nice properties. First, q is convex. Second, it grows asymptotically log-linearly. Third, while $\exp(-|\log t|) = \exp(-\max(t, 1/t))$ is a seemingly natural choice for normalization, it is non smooth, and the LogSumExp [12] replacement of $\max(t, 1/t)$ with $\log(\exp(t) + \exp(1/t))$ simplifies to 27 upon exponentiation. This normalization provides good numerical properties for avoiding issues arising close to 0 and ∞. We plot the effect of this normalization on our example data in Figure 1c.

3.3 Isometry pursuit

Isometry pursuit is the application of multitask basis pursuit to the normalized design matrix $w(X)$ to identify submatrices of X that are as orthonormal as possible. Define the multitask basis pursuit penalty

$$\|\cdot\|_{1,2} : \mathbb{R}^{P\times D} \to \mathbb{R}^+ \tag{28}$$

$$\beta \mapsto \sum_{p=1}^{P} \|\beta_{p\cdot}\|_2. \tag{29}$$

Given a matrix $Y \in \mathbb{R}^{D\times D}$, the multitask basis pursuit solutions are

$$\widehat{\beta}_{MBP}(X,Y) := \arg \min_{\beta\in\mathbb{R}^{P\times D}} \|\beta\|_{1,2} \ : \ Y = X\beta. \tag{30}$$

The multitask basis pursuit solution is often not unique. See Schneider and Tardivel [22] for more details on uniqueness of basis pursuit solutions, for which the uniqueness condition is less strict than in the multitask case. We therefore define

$$\widehat{\beta}^{\ell_2}_{MLP}(X,Y) := \arg\min \|\beta\|_F : \beta \in \beta_{MBP}(X,Y). \tag{31}$$

By the strong convexity of the ℓ_2 norm, $\widehat{\beta}^{\ell_2}_{MLP}(X,Y)$ is well defined.

This implicit regularization assumption is reasonable. For example, in Tibshirani [23] the implicit selection of the minimum ℓ_2 norm solution among the lasso solutions was proven for the Least Angle Regression Selection (LARS) algorithm for solving the lasso, while Mishkin and Pilanci [24] assumes that the group lasso solution is in fact min-norm prior to subsequent theoretical analyses. Empirically, the Splitting Conic Solver type method (which is itself a variant of the Alternating Direction Method of Multilpliers) [25] stably estimates the min-norm solution when initialized at 0, but besides a brief exploration in Section **??** we leave theoretical and experimental proof of this feature of the feature our optimization approach for future work, and instead make the ability of our optimizer to obtain the min-norm solution an assumption of our eventual proposition.

Isometry pursuit is then given by

$$\widehat{\beta}(X) := \widehat{\beta}^{\ell_2}_{MBP}(w(X), I_D) \tag{32}$$

where I_D is the D dimensional identity matrix and recovered functions are the indices of the dictionary elements with non-zero coefficients. That is, they are given by $S(\beta)$ where

$$S : \mathbb{R}^{P\times D} \to \binom{[P]}{D} \tag{33}$$

$$\beta \mapsto \{p \in [P] : \|\beta_{p\cdot}\| > 0\}. \tag{34}$$

IsometryPursuit(Matrix $X \in \mathbb{R}^{D\times P}$)

1: Normalize $X' = w(X)$
2: Optimize $\widehat{\beta} = \widehat{\beta}^{\ell_2}_{MBP}(X', I_D)$
3: **Output** $\widehat{S} = S(\widehat{\beta})$

3.4 Theory

The intuition behind our application of multitask basis pursuit is that submatrices consisting of vectors which are closer to 1 in length and more orthogonal will have smaller loss. A key theoretical assertion is that IsometryPursuit is invariant to choice of basis for X.

Proposition 9 *Let $U \in \mathbb{R}^{D\times D}$ be orthonormal. Then $S(\widehat{\beta}(UX)) = S(\widehat{\beta}(X))$.*

This proposition establishes the basis-independence of Isometry Pursuit and that we may replace I_D in the constraint by any orthonormal $D \times D$ matrix. It is also essential for showing the following main theoretical result.

When a rank D orthonormal column-submatrix $X_{\cdot S}$ exists, the output of Program 32 will contain S.

Proposition 10 *Let $X \in \mathbb{R}^{D\times P}$ have a rank D orthonormal column submatrix $X_{\cdot S}$. Then $S \subseteq \widehat{\beta}(X)$.*

Proofs of these propositions are given in Section 6.2, with corresponding experimental results in Section 4. The proof mostly relies on the KKT conditions [26] being used to verify that an isometric subset yields a corresponding basis pursuit solution for an appropriate normalized matrix. This is a formalization of our intuition that longer and more orthogonal vectors will minimize loss. We then apply a simple argument to show that the ℓ_2 minimizing solution is the least sparse solution, which means that it contains the isometry solution, should it exist.

3.5 Two-stage isometry pursuit

Proposition 10 suggests the following algorithm, which first uses the convex IsometryPursuit algorithm to prune the P candidate features and then applies brute search upon the reduced feature set. The proposition says that this approach is guaranteed to select an isometric submatrix should one exist. The two-stage approach is standard in the Lasso literature [27, 28, 10], and Proposition 10 justifies it in our application. This method forms our practical isometry estimator, the core advantage of which over purely brute search is a substantially improved computational complexity.

TwoStageIsometryPursuit(Matrix $X \in \mathbb{R}^{D \times P}$)

1: $\widehat{S}_{IP} = \text{IsometryPursuit}(X)$
2: $\widehat{S} = \text{BruteSearch}(X_{\cdot \widehat{S}_{IP}}, \ell)$
3: **Output** $\widehat{S}$

3.6 Non-uniqueness

The isometric solution is optimum among cardinality D solutions.

Proposition 11 $\ell(X_S) = D, |S| = D \Leftrightarrow X_S X_S^T = I_D$

However, there exist Isometry Pursuit solutions of minimum norm that are not isometries. For example, *Mercedes-Benz frame* matrix X achieves minimium norm.

Example 1 *Let* $D = 2, P = 3$. *Let* $X_S = \begin{bmatrix} 0 & \frac{\sqrt{3}}{2} & -\frac{\sqrt{3}}{2} \\ 1 & -\frac{1}{2} & -\frac{1}{2} \end{bmatrix}, \beta = \begin{bmatrix} 0 & \frac{2}{3} \\ \frac{1}{\sqrt{3}} & -\frac{1}{3} \\ -\frac{1}{\sqrt{3}} & -\frac{1}{3} \end{bmatrix}$ *Then* $X\beta = I_D$ *and* $\|\beta\|_{1,2} = \frac{2}{3} + \frac{2}{3} + \frac{2}{3} = 2 = D$. *Since the identity matrix is a solution regardless of* X_{-S}, *this is a non-isometric minimum loss solution.*

The upshot of this example is that IsometryPursuit can select solutions that are relatively non-isometric. The Mercedes-Benz frame is an example of a tight frame - a frame whose non-zero singular values are equal. While there are non-tight frames that have minimizing solutions, we do not here resolve whether there exist tight frames with unit norm columns that have minimum norm solutions.

3.7 Computation complexity

IsometryPursuit is more computationally efficient than BruteSearch search for the minimizer of ℓ. Interestingly, their asymptotic per-iteration complexities are the same.

Proposition 12 *The computational complexity* C_ℓ *of 10 is* $O(D^2 P)$.

The above proposition follows immediately from the computational complexity of singular value decomposition [29]. The per-iteration complexity of IsometryPursuit is also limited by the computational complexity of singular value decomposition, albeit in the more intricate manner described in Section 6.2.6.

Proposition 13 *The computational complexity of each step of Isometry Pursuit is* $O(D^2 P)$.

Figure 5: Difference in isometry losses $\ell(\widehat{S}_g) - \ell(\widehat{S}_{TSIP})$ for Wine, Iris, and Ethanol datasets. The distributions reflect the relative losses of the supports estimated by greedy optimization on the ℓ objective is less than or equal to ℓ loss of the support estimated by TwoStageIsometryPursuit. As shown here and detailed in Table 3, losses are generally lower for two-stage isometry pursuit solutions.

Name	D	P	R	$\ell(X_{\cdot\widehat{S}_G})$	$\lvert\widehat{S}\rvert$	$\ell(X_{\cdot\widehat{S}})$	$\widehat{P}(\ell(X_{\cdot\widehat{S}_G}) > \ell(X_{\cdot\widehat{S}}))$	$\widehat{P}(\ell(X_{\cdot\widehat{S}_G}) = \ell(X_{\cdot\widehat{S}}))$	$\widehat{P}(\bar{\ell}_1(X_{\cdot\widehat{S}_G}) > \bar{\ell}_1(X_{\cdot\widehat{S}}))$
Iris	4	75	25	14 ± 7.3	6.7 ± 1	6.9 ± 1.4	0.96	0	2.4e-05
Wine	6	89	25	7.7 ± 0.33	13 ± 1.5	7.6 ± 0.29	0.64	0.16	6.3e-04
Ethanol	2	756	25	2.6 ± 0.3	90 ± 1.7e+02	2.5 ± 0.2	0.66	0.17	2.1e-05

Table 3: Experimental parameters and results. For Iris and Wine, probability estimates $\widehat{P}$ result from random downsampling features $[P]$ by a factor of 2 to create R replicates, while for Ethanol, replicates correspond to individual data points. Distributional probabilities $\widehat{P}(\ell(X_{\cdot\widehat{S}_G}) > \ell(X_{\cdot\widehat{S}_{TSIP}}))$ and $\widehat{P}(\ell(X_{\cdot\widehat{S}_G}) = \ell(X_{\cdot\widehat{S}_{TSIP}}))$ are empirical across replicates, while asymptotic probabilities $\widehat{P}(\bar{\ell}_1(X_{\cdot\widehat{S}_G}) > \bar{\ell}_1(X_{\cdot\widehat{S}_{TSIP}}))$ is computed by paired two-sample T-test on $\ell(X_{\cdot\widehat{S}})$ and $\ell(X_{\cdot\widehat{S}_G})$. For brevity, in this table $\widehat{S} := \widehat{S}_{TSIP}$.

The difference therefore lies in the relation of the number of iterations required to achieve a given error tolerance with IsometryPursuit and BruteSearch. To optimize, we use the SCS interior point solver [30] from CVXPY [31, 32], which is able to push sparse values arbitrarily close to 0 [33]. SCS (aka ADMM) has asymptotic (in error) accuracy $\epsilon < O(\frac{1}{\mathcal{S}})$, where $\mathcal{S}$ is the number of iterations [34]. This means that the number of iterations required to achieve a given error tolerance scales inversely with the desired level of accuracy. We examine the effect of iteration number on the convergence to a sparse solution in Section 4.

4 Experiments

Say you are hosting a dinner party, and wish to select a balanced set of wines for drinking and flowers for decoration. We demonstrate TwoStageIsometryPursuit and GreedySearch on the Iris and Wine datasets [35, 36, 37]. This has an intuitive interpretation as selecting diverse elements that reflects the peculiar structure of the diversification problem. Features like *petal width* are rows in X. Thus, in diversification, $P = n$. Due to differences in scales between features, these are first standardized. For the Wine dataset, even BruteSearch on $\widehat{S}_{IP}$ is prohibitive in $D = 13$, and so we sample our inputs to $D = 6$.

Our main interpretability dataset is the Ethanol dataset from Chmiela et al. [38], Koelle et al. [10]. For this dataset evaluate a dictionary of interpretable features - bond torsions - for their ability to parameterize the molecular configuration space. In this interpretability use case, columns denote gradients of informative features. We compute Jacoban matrices of putative parametrization functions and project them onto estimated tangent spaces (see Koelle et al. [10] for preprocessing details). Rather than selecting between data points, we are selecting between functions which parameterize the data. Replicas are created by sampling from data points and their corresponding tangent spaces are estimated in $B = 252$. This dataset exemplifies our motivating example - the search for locally isometric embeddings.

Figure 5 and Table 3 show that the ℓ accrued by the subset $\widehat{S}_G$ estimated using GreedySearch with objective ℓ is higher than that for the subset estimated by TwoStageIsometryPursuit. This effect is statistically significant across datasets. Generally, the ground truth loss obtained by GreedySearch optimization of ℓ is higher and therefore worse than that obtained by two-stage isometry pursuit. Table 3 and Figure 6 detail intermediate support recovery cardinalities from IsometryPursuit. Isometry pursuit substantially reduced the size of the dictionary and made brute search computationally feasible. Wall-clock runtimes are given in Section 6.5.

5 Discussion

We have shown that multitask basis pursuit can help select isometric submatrices from appropriately normalized wide matrices. This approach - IsometryPursuit - is a convex alternative to greedy methods for selection of orthonormalized features from within a dictionary. Isometry pursuit can be applied to diversification and geometrically-faithful coordinate estimation. Our results exemplify these applications, but more can be done.

On the theoretical side, a critical unresolved question is the relation of the brute force and IsometryPursuit minimizers. Essentially, there are two things that can go wrong. First, IsometryPursuit may trim insufficiently for brute force search to be computationally feasible. Second, the recovered support may not contain the true minimizer. While we have ruled out the second problem in the case where an isometric solution does exist, in practice we often see greedy recovery defeat IsometryPursuit. Characterization of support recovery here is somewhat different from existing analysis of greedy [39, 40, 41, 42] and convex [18, 43, 44] basis pursuit formulations. Partly, this is because we use the ℓ_1 minimizer to distinguish between ℓ_0 optima rather than to identify them. Also, IsometryPursuit can generate non-unique solutions even when the design matrix satisfies general position, which is not true for the univariate formulation [45, 46, 47, 48, 49, 50, 51, 52, 53, 54, 55], but

The principal algorithmic variant of interest is the combination of the multitask lasso [26] extension of our estimator with D function selection within $\mathbb{R}^B$. First, introducing a Lasso penalty may ameliorate IsometryPursuit's reliance on the second stage of coordinate selection to complete the estimation process. Using the Lasso penalty, we can smoothly move between the multitask basis pursuit solution and more sparse solutions [24]. Second, we often may want to select some orthonormal subset of dimension much less than D. Finally, response matrices may be extended beyond the identity.

These extensions may faciliate an interesting set of applications diversification in recommendation [56, 57, 58] and retrieval [59, 60, 61, 62, 63]. Maximum cosine similarity used in retrieval corresponds to the first coefficient of the Lasso regularization path [10]. Greedy algorithms predominate in diversification [64, 65, 66, 67, 68, 69, 70, 71, 72, 60, 2, 73], and convex reformulations have been studied to add speed and convergence to a global minima.

Compared with Sparse PCA [74, 75, 76, 77], we seek a low-dimensional representation from the set of these column vectors rather than their span.

Diversification.

While other applications like diversification are possible and potentially , the embedding

Diversification may be framed an a kind of orthogonal selection problem. Cosine similarity Colbert

References

[1] Samson J Koelle, Hanyu Zhang, Octavian-Vlad Murad, and Marina Meila. Consistency of dictionary-based manifold learning. In Sanjoy Dasgupta, Stephan Mandt, and Yingzhen Li, editors, *Proceedings of The 27th International Conference on Artificial Intelligence and Statistics*, volume 238 of *Proceedings of Machine Learning Research*, pages 4348–4356. PMLR, 2024.

[2] T. Tony Cai and Lie Wang. Orthogonal matching pursuit for sparse signal recovery with noise. *IEEE Transactions on Information Theory*, 57(7):4680–4688, 2011. doi: 10.1109/TIT.2011.2146090.

[3] Yu-Chia Chen and Marina Meila. Selecting the independent coordinates of manifolds with large aspect ratios. In H. Wallach, H. Larochelle, A. Beygelzimer, F. d'Alché-Buc, E. Fox, and R. Garnett, editors, *Advances in Neural Information Processing Systems*, volume 32. Curran Associates, Inc., 2019. URL `https://proceedings.neurips.cc/paper_files/paper/2019/file/6a10bbd480e4c5573d8f3af73ae0454b-Paper.pdf`.

[4] Dhruv Kohli, Alexander Cloninger, and Gal Mishne. LDLE: Low distortion local eigenmaps. *J. Mach. Learn. Res.*, 22, 2021.

[5] Peter W Jones, Mauro Maggioni, and Raanan Schul. Universal local parametrizations via heat kernels and eigenfunctions of the laplacian. September 2007.

[6] Zhenyue Zhang and Hongyuan Zha. Principal manifolds and nonlinear dimensionality reduction via tangent space alignment. *SIAM J. Scientific Computing*, 26(1):313–338, 2004.

[7] Lisha Chen and Andreas Buja. Local Multidimensional Scaling for nonlinear dimension reduction, graph drawing and proximity analysis. *Journal of the American Statistical Association*, 104(485):209–219, March 2009.

[8] J.B. Tenenbaum, V. Silva, and J.C. Langford. A global geometric framework for nonlinear dimensionality reduction. *Science*, 290(5500):2319–2323, 2000.

[9] John M. Lee. *Introduction to Riemannian Manifolds*, volume 176 of *Graduate Texts in Mathematics*. Springer, Cham, 2 edition, 2018. doi: 10.1007/978-3-319-91755-9.

[10] Samson J Koelle, Hanyu Zhang, Marina Meila, and Yu-Chia Chen. Manifold coordinates with physical meaning. *J. Mach. Learn. Res.*, 23(133):1–57, 2022.

[11] Gene H. Golub and James H. Wilkinson. Condition number. URL `https://gdz.sub.uni-goettingen.de/id/PPN362160546_0009`.

[12] Stephen P Boyd and Lieven Vandenberghe. *Convex Optimization*. Cambridge University Press, March 2004.

[13] Suvrit Sra. Positive definite matrices and the S-divergence, 2011. URL `https://arxiv.org/pdf/1110.1773`.

[14] Andrzej Cichocki, Sergio Cruces, and Shun-ichi Amari. Log-determinant divergences revisited: Alpha-beta and gamma log-det divergences, 2014. URL `https://arxiv.org/pdf/1412.7146`.

[15] Leo Sergio Liberti. *Reformulation and Convex Relation Techniques for Global Optimization*. PhD thesis, Imperial College London, October 2004.

[16] M Abdi. *Cardinality optimization problems*. PhD thesis, December 2013.

[17] Fengyu Zhou and Steven H Low. Conditions for exact convex relaxation and no spurious local optima. *arXiv [math.OC]*, February 2021.

[18] Scott Shaobing Chen and David L. Donoho and Michael A. Saunders. Atomic decomposition by basis pursuit. *SIAM REVIEW*, 43(1):129, February 2001.

[19] Trevor Hastie, Robert Tibshirani, and Ryan Tibshirani. Best subset, forward stepwise, or lasso?, 2016.

[20] Ulrike Schneider and Patrick Tardivel. The geometry of uniqueness, sparsity and clustering in penalized estimation. *Journal of Machine Learning Research*, 23(331):1–36, 2022. URL `http://jmlr.org/papers/v23/21-0420.html`.

[21] Dávid Szeghy, Mahmoud Aslan, Áron Fóthi, Balázs Mészáros, Zoltán Ádám Milacski, and András Lőrincz. Structural extensions of basis pursuit: Guarantees on adversarial robustness. *arXiv [cs.LG]*, May 2022.

[22] Ulrike Schneider and Patrick Tardivel. The geometry of uniqueness, sparsity and clustering in penalized estimation. *J. Mach. Learn. Res.*, 23(331):1–36, 2022.

[23] Ryan J Tibshirani. The lasso problem and uniqueness. *arXiv [math.ST]*, June 2012.

[24] Aaron Mishkin and Mert Pilanci. The solution path of the group lasso. 2022.

[25] Brendan O'Donoghue, Eric Chu, Neal Parikh, and Stephen Boyd. Conic optimization via operator splitting and homogeneous self-dual embedding. *arXiv [math.OC]*, December 2013.

[26] T Hastie, R Tibshirani, and M Wainwright. Statistical learning with sparsity: The lasso and generalizations. May 2015.

[27] Nicolai Meinshausen. Relaxed lasso. *Comput. Stat. Data Anal.*, 52(1):374–393, September 2007.

[28] Tim Hesterberg, Nam Hee Choi, Lukas Meier, and Chris Fraley. Least angle and ℓ_1 penalized regression: A review. February 2008.

[29] Faming Liang, Runmin Shi, and Qianxing Mo. A split-and-merge approach for singular value decomposition of large-scale matrices. *Stat. Interface*, 9(4):453–459, September 2016.

[30] Brendan O'Donoghue, Eric Chu, Neal Parikh, and Stephen Boyd. Conic optimization via operator splitting and homogeneous self-dual embedding. *Journal of Optimization Theory and Applications*, 169(3):1042–1068, June 2016. URL `http://stanford.edu/~boyd/papers/scs.html`.

[31] Steven Diamond and Stephen Boyd. CVXPY: A Python-embedded modeling language for convex optimization. *Journal of Machine Learning Research*, 17(83):1–5, 2016.

[32] Akshay Agrawal, Robin Verschueren, Steven Diamond, and Stephen Boyd. A rewriting system for convex optimization problems. *Journal of Control and Decision*, 5(1):42–60, 2018.

[33] CVXPY Developers. Sparse solution with cvxpy. `https://www.cvxpy.org/examples/applications/sparse_solution.html`. Accessed: 2024-07-11.

[34] Bingsheng He and Xiaoming Yuan. On the $o(1/n)$ convergence rate of the douglas–rachford alternating direction method. *SIAM J. Numer. Anal.*, 50(2):700–709, January 2012.

[35] R. A. Fisher. Iris. UCI Machine Learning Repository, 1988. DOI: https://doi.org/10.24432/C56C76.

[36] Stefan Aeberhard and M. Forina. Wine. UCI Machine Learning Repository, 1991. DOI: https://doi.org/10.24432/C5PC7J.

[37] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and E. Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.

[38] Stefan Chmiela, Huziel E Sauceda, Klaus-Robert Müller, and Alexandre Tkatchenko. Towards exact molecular dynamics simulations with machine-learned force fields. *Nat. Commun.*, 9(1):3887, September 2018.

[39] S.G. Mallat and Zhifeng Zhang. Matching pursuits with time-frequency dictionaries. *IEEE Transactions on Signal Processing*, 41(12):3397–3415, 1993. doi: 10.1109/78.258082.

[40] S. Mallat and Z. Zhang. Adaptive time-frequency decomposition with matching pursuits. In *[1992] Proceedings of the IEEE-SP International Symposium on Time-Frequency and Time-Scale Analysis*, pages 7–10, 1992. doi: 10.1109/TFTSA.1992.274245.

[41] Y.C. Pati, R. Rezaiifar, and P.S. Krishnaprasad. Orthogonal matching pursuit: recursive function approximation with applications to wavelet decomposition. In *Proceedings of 27th Asilomar Conference on Signals, Systems and Computers*, pages 40–44 vol.1, 1993. doi: 10.1109/ACSSC.1993.342465.

[42] J.A. Tropp, A.C. Gilbert, and M.J. Strauss. Simultaneous sparse approximation via greedy pursuit. In *Proceedings. (ICASSP '05). IEEE International Conference on Acoustics, Speech, and Signal Processing, 2005.*, volume 5, pages v/721–v/724 Vol. 5, 2005. doi: 10.1109/ICASSP.2005.1416405.

[43] Joel A. Tropp. Algorithms for simultaneous sparse approximation. part ii: Convex relaxation. *Signal Processing*, 86(3):589–602, 2006. ISSN 0165-1684. doi: https://doi.org/10.1016/j.sigpro.2005.05.031. URL https://www.sciencedirect.com/science/article/pii/S0165168405002239. Sparse Approximations in Signal and Image Processing.

[44] Jie Chen and Xiaoming Huo. Theoretical results on sparse representations of multiple-measurement vectors. *IEEE Transactions on Signal Processing*, 54:4634–4643, 2006. URL https://api.semanticscholar.org/CorpusID:17333301.

[45] Michael R. Osborne, Brett Presnell, and Berwin A. Turlach. On the lasso and its dual. *Journal of Computational and Graphical Statistics*, 9:319 – 337, 2000. URL https://api.semanticscholar.org/CorpusID:14422381.

[46] Charles Dossal. A necessary and sufficient condition for exact sparse recovery by l1 minimization. *Comptes Rendus Mathematique*, 350(1):117–120, 2012. ISSN 1631-073X. doi: https://doi.org/10.1016/j.crma.2011.12.014. URL https://www.sciencedirect.com/science/article/pii/S1631073X11003694.

[47] Stéphane Chrétien and Sébastien Darses. On the generic uniform uniqueness of the lasso estimator. *arXiv: Statistics Theory*, 2011. URL https://api.semanticscholar.org/CorpusID:88518316.

[48] Ryan J. Tibshirani. The lasso problem and uniqueness. *Electronic Journal of Statistics*, 7:1456–1490, 2012. URL https://api.semanticscholar.org/CorpusID:5849668.

[49] Karl Ewald and Ulrike Schneider. On the distribution, model selection properties and uniqueness of the lasso estimator in low and high dimensions. *Electronic Journal of Statistics*, 2017. URL https://api.semanticscholar.org/CorpusID:54044415.

[50] Alnur Ali and Ryan J. Tibshirani. The generalized lasso problem and uniqueness. *Electronic Journal of Statistics*, 2018. URL https://api.semanticscholar.org/CorpusID:51755233.

[51] Ulrike Schneider and Patrick Tardivel. The geometry of uniqueness, sparsity and clustering in penalized estimation. *arXiv [math.ST]*, April 2020.

[52] Aaron Mishkin and Mert Pilanci. The solution path of the group lasso. 2022. URL https://api.semanticscholar.org/CorpusID:259504228.

[53] Xavier Dupuis and Samuel Vaiter. The geometry of sparse analysis regularization. *SIAM J. Optim.*, 33:842–867, 2019. URL https://api.semanticscholar.org/CorpusID:195791526.

[54] Thomas Debarre, Quentin Denoyelle, and Julien Fageot. On the uniqueness of solutions for the basis pursuit in the continuum. *Inverse Problems*, 38, 2020. URL https://api.semanticscholar.org/CorpusID:246473440.

[55] Jasper Marijn Everink, Yiqiu Dong, and Martin Skovgaard Andersen. The geometry and well-posedness of sparse regularized linear regression. 2024. URL https://api.semanticscholar.org/CorpusID:272424099.

[56] Jaime Carbonell and Jade Goldstein. The use of MMR, diversity-based reranking for reordering documents and producing summaries. *SIGIR Forum*, 51(2):209–210, August 1998.

[57] Qiong Wu, Yong Liu, Chunyan Miao, Yin Zhao, Lu Guan, and Haihong Tang. Recent advances in diversified recommendation. May 2019.

[58] Select by maximal marginal relevance (MMR). https://python.langchain.com/docs/how_to/example_selectors_mmr/. Accessed: 2024-11-22.

[59] Yunfan Gao, Yun Xiong, Xinyu Gao, Kangxiang Jia, Jinliu Pan, Yuxi Bi, Yi Dai, Jiawei Sun, Meng Wang, and Haofen Wang. Retrieval-augmented generation for large language models: A survey. *arXiv [cs.CL]*, December 2023.

[60] Marc Pickett, Jeremy Hartman, Ayan Kumar Bhowmick, Raquib-Ul Alam, and Aditya Vempaty. Better RAG using relevant information gain. *arXiv [cs.CL]*, July 2024.

[61] Yeonjun In, Sungchul Kim, Ryan A Rossi, Md Mehrab Tanjim, Tong Yu, Ritwik Sinha, and Chanyoung Park. Diversify-verify-adapt: Efficient and robust retrieval-augmented ambiguous question answering. *arXiv [cs.CL]*, September 2024.

[62] Sam Weiss. Enhancing diversity in RAG document retrieval using projection-based techniques. https://medium.com/@samcarlos_14058/enhancing-diversity-in-rag-document-retrieval-using-projection-based-techniques-9fef5422e043 August 2024. Accessed: 2024-11-22.

[63] Get diverse results and comprehensive summaries with vectara's MMR reranker. URL https://www.vectara.com/blog/get-diverse-results-and-comprehensive-summaries-with-vectaras-mmr-reranker. Accessed: 2024-11-22.

[64] Jaime Carbonell and Jade Goldstein. The use of MMR, diversity-based reranking for reordering documents and producing summaries. In *Proceedings of the 21st annual international ACM SIGIR conference on Research and development in information retrieval*, New York, NY, USA, August 1998. ACM.

[65] Maria C. N. Barioni, Marios Hadjieleftheriou, Marcos R. Vieira, Caetano Traina, Vassilis J. Tsotras, Humberto L. Razente, and Divesh Srivastava. On query result diversification . In *2011 27th IEEE International Conference on Data Engineering (ICDE 2011)*, pages 1163–1174, Los Alamitos, CA, USA, April 2011. IEEE Computer Society. doi: 10.1109/ICDE.2011.5767846. URL https://doi.ieeecomputersociety.org/10.1109/ICDE.2011.5767846.

[66] Marina Drosou and Evaggelia Pitoura. Search result diversification. *SIGMOD Rec.*, 39(1):41–47, September 2010. ISSN 0163-5808. doi: 10.1145/1860702.1860709. URL https://doi.org/10.1145/1860702.1860709.

[67] Lu Qin, Jeffrey Xu Yu, and Lijun Chang. Diversifying top-k results. *Proceedings VLDB Endowment*, 5(11):1124–1135, July 2012.

[68] Matevž Kunaver and Tomaž Požrl. Diversity in recommender systems – a survey. *Knowledge-Based Systems*, 123:154–162, 2017. ISSN 0950-7051. doi: https://doi.org/10.1016/j.knosys.2017.02.009. URL https://www.sciencedirect.com/science/article/pii/S0950705117300680.

[69] Shengbo Guo and Scott Sanner. Probabilistic latent maximal marginal relevance. In *Proceedings of the 33rd International ACM SIGIR Conference on Research and Development in Information Retrieval*, SIGIR '10, page 833–834, New York, NY, USA, 2010. Association for Computing Machinery. ISBN 9781450301534. doi: 10.1145/1835449.1835639. URL https://doi.org/10.1145/1835449.1835639.

[70] Mustafa Abdool, Malay Haldar, Prashant Ramanathan, Tyler Sax, Lanbo Zhang, Aamir Manaswala, Lynn Yang, Bradley Turnbull, Qing Zhang, and Thomas Legrand. Managing diversity in airbnb search. In *Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, KDD '20, page 2952–2960, New York, NY, USA, 2020. Association for Computing Machinery. ISBN 9781450379984. doi: 10.1145/3394486.3403345. URL https://doi.org/10.1145/3394486.3403345.

[71] Hsiang-Fu Yu, Cho-Jui Hsieh, Qi Lei, and Inderjit S. Dhillon. A greedy approach for budgeted maximum inner product search. In *Neural Information Processing Systems*, 2016. URL https://api.semanticscholar.org/CorpusID:7076785.

[72] Qiang Huang, Yanhao Wang, Yiqun Sun, and Anthony K H Tung. Diversity-aware k-maximum inner product search revisited. *arXiv [cs.IR]*, February 2024.

[73] Jianchen Zhu, Shengjie Zhao, and Rongqing Zhang. A match pursuit based method adapted to overcomplete dictionary for compressive spectral imaging. In *2020 Asia-Pacific Signal and Information Processing Association Annual Summit and Conference (APSIPA ASC)*, pages 170–173. IEEE, December 2020.

[74] Alexandre d'Aspremont, Laurent El Ghaoui, Michael I Jordan, and Gert R G Lanckriet. A direct formulation for sparse PCA using semidefinite programming. *SIAM Rev. Soc. Ind. Appl. Math.*, 49(3):434–448, January 2007.

[75] Santanu S Dey, R Mazumder, M Molinaro, and Guanyi Wang. Sparse principal component analysis and its ℓ_1-relaxation. *arXiv: Optimization and Control*, December 2017.

[76] D Bertsimas and Driss Lahlou Kitane. Sparse PCA: A geometric approach. *J. Mach. Learn. Res.*, 24:32:1–32:33, October 2022.

[77] Dimitris Bertsimas, Ryan Cory-Wright, and Jean Pauphilet. Solving Large-Scale sparse PCA to certifiable (near) optimality. *J. Mach. Learn. Res.*, 23(13):1–35, 2022.

[78] Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein. *Introduction to Algorithms, Third Edition*. The MIT Press, 3rd edition, 2009. ISBN 0262033844.

[79] Stuart Russell and Peter Norvig. *Artificial Intelligence: A Modern Approach*. Prentice Hall Press, USA, 3rd edition, 2009. ISBN 0136042597.

[80] Travis Dick, Eric Wong, and Christoph Dann. How many random restarts are enough. 2014. URL https://api.semanticscholar.org/CorpusID:9473630.

[81] Stephen Boyd. Distributed optimization and statistical learning via the alternating direction method of multipliers. *Found. Trends® Mach. Learn.*, 3(1):1–122, 2010.

6 Supplement

This section contains algorithms, proofs, and experiments in support of the main text.

6.1 Algorithms

We give definitions of the brute and greedy algorithms for the combinatorial problem studied in this paper. The brute force algorithm is computationally intractable for all but the smallest problems, but always finds the global minima.

BruteSearch(Matrix $X \in \mathbb{R}^{D \times P}$, objective f)

1: **for** each combination $S \subseteq \{1, 2, \ldots, P\}$ with $|S| = D$ **do**
2: Evaluate $f(X_{.S})$
3: **end for**
4: **Output** the combination S^* that minimizes $f(X_{.S})$

Greedy algorithms are computationally expedient but can get stuck in local optima [78, 79], even with randomized restarts [80].

GreedySearch(Matrix $X \in \mathbb{R}^{D \times P}$, objective f, selected set $S = \emptyset$, current size $d = 0$)

1: **if** $d = D$ **then**
2: **Return** S
3: **else**
4: **Initialize** $S_{\text{best}} = S$
5: **Initialize** $f_{\text{best}} = \infty$
6: **for** each $p \in \{1, 2, \ldots, P\} \setminus S$ **do**
7: **Evaluate** $f(X_{.(S \cup \{p\})})$
8: **if** $f(X_{.(S \cup \{p\})}) < f_{\text{best}}$ **then**
9: **Update** $S_{\text{best}} = S \cup \{p\}$
10: **Update** $f_{\text{best}} = f(\mathcal{X}_{.(S \cup \{p\})})$
11: **end if**
12: **end for**
13: **Return** GreedySearch(X, f, S_{best}, $d + 1$)
14: **end if**

6.2 Proofs

6.2.1 Proof of Proposition 6

Proof: First,

$$\frac{1}{2} \log \det(X^T X) = \frac{1}{2} \log \prod_{d=1}^{D} \sigma_d(X)^2 \tag{35}$$

$$= \log \prod_{d=1}^{D} \sigma_d(X) \tag{36}$$

$$= \sum_{d=1}^{D} \log \sigma_d(X). \tag{37}$$

Secondly, I_D and X^TX are diagonalizable in the same basis, so

$$\log\det(\frac{X^TX+I_D}{2}) = \log\prod_{d=1}^{D}\frac{\sigma_d(X)^2+1}{2} \tag{38}$$

$$= \sum_{d=1}^{D}\log\frac{\sigma_d(X)^2+1}{2}. \tag{39}$$

Then,

$$D(X^TX, I_D)^2 = \sum_{d=1}^{D}(\log\frac{\sigma_d(X)^2+1}{2} - \log\sigma_d(X)) \tag{40}$$

$$= \sum_{d=1}^{D}\log\frac{\sigma_d(X)^2+1}{2\sigma_d(X)} \tag{41}$$

$$= \sum_{d=1}^{D}\log\frac{\sigma_d(X)+\sigma_d(X)^{-1}}{2}. \tag{42}$$

□

6.2.2 Singular values of nearly colinear vectors

We now exhibit a matrix for which no g can be chosen such that a log det divergence satisfies Properties **??** and **??**. We parametrize this matrix as

$$X_\rho(r) = \begin{bmatrix} r & \rho & \rho \\ 0 & \sqrt{1-\rho^2} & \frac{\rho(1-\rho)}{\sqrt{1-\rho^2}} \\ 0 & 0 & \sqrt{\frac{(1-\rho)(1+2\rho)}{1+\rho}} \end{bmatrix} \tag{43}$$

where r is the length of the second vector and ρ controls the colinearity between vectors. We compute

$$\frac{\partial\ell(r)s}{\partial r}|_{r=1} = \sum_{d=1}^{3} g'(\sigma_d(X_\rho(r)))\frac{\partial}{\partial r}\sigma_d(X_\rho(r))\Big|_{r=1}. \tag{44}$$

We will to show that, for any log-symmetric g such that $g(t) = g(1/t)$ and any $\rho \in (-1, 1)$

$$\frac{\partial\ell(r)}{\partial r}|_{r=1} \neq 0. \tag{45}$$

First, let's compute $\sigma_d(X_\rho(r))$.

$$X_\rho(r)^TX_\rho(r) = \begin{bmatrix} r^2 & r\rho & r\rho \\ r\rho & 1 & \rho \\ r\rho & \rho & 1 \end{bmatrix}. \tag{46}$$

Recall that λ is an eigenvalue of a matrix A if there is a non-zero vector v such that $Av = \lambda v$. Equivalently,

$$\det(A - \lambda I) = 0. \tag{47}$$

In our case, this gives

$$0 = \det \begin{bmatrix} r^2-\lambda & r\rho & r\rho \\ r\rho & 1-\lambda & \rho \\ r\rho & \rho & 1-\lambda \end{bmatrix} \tag{48}$$

$$= (r^2-\lambda)\left((1-\lambda)^2-\rho^2\right) - r\rho\left(r\rho(1-\lambda) - r\rho^2\right) + r\rho\left(r\rho^2 - r\rho(1-\lambda)\right) \tag{49}$$

$$= (r^2-\lambda)\left((1-\lambda)^2-\rho^2\right) - 2r^2\rho^2(1-\lambda) + 2r^2\rho^3 \tag{50}$$

$$= (r^2-\lambda)(1-\rho-\lambda)(1+\rho-\lambda) - 2r^2\rho^2(1-\rho-\lambda) \tag{51}$$

$$= (1-\rho-\lambda)\left((r^2-\lambda)(1+\rho-\lambda) - 2r^2\rho^2\right) \tag{52}$$

$$= (1-\rho-\lambda)\left(r^2(1+\rho) - r^2\lambda - \lambda(1+\rho) + \lambda^2 - 2r^2\rho^2\right) \tag{53}$$

$$= (1-\rho-\lambda)\left(\lambda^2 - (r^2+1+\rho)\lambda + r^2(1+\rho-2\rho^2)\right). \tag{54}$$

Thus the eigenvalues of $X_\rho(r)^T X_\rho(r)$ are

$$\lambda_1(r) = 1-\rho, \tag{55}$$

$$\lambda_2(r) = \frac{r^2+1+\rho+\sqrt{(r^2+1+\rho)^2 - 4r^2(1+\rho-2\rho^2)}}{2} \tag{56}$$

$$= \frac{r^2+1+\rho+\sqrt{\left(r^2-(1+\rho)\right)^2 + 8r^2\rho^2}}{2}, \tag{57}$$

$$\lambda_3(r) = \frac{r^2+1+\rho-\sqrt{(r^2+1+\rho)^2 - 4r^2(1+\rho-2\rho^2)}}{2} \tag{58}$$

$$= \frac{r^2+1+\rho-\sqrt{\left(r^2-(1+\rho)\right)^2 + 8r^2\rho^2}}{2}. \tag{59}$$

Since singular values are the square roots of the eigenvalues of $X_\epsilon(r)^T X_\epsilon(r)$,

$$\sigma_1(X_\rho(r)) = \left[1-\rho\right]^{1/2}, \tag{60}$$

$$\sigma_2(X_\rho(r)) = \left[\frac{r^2+1+\rho+\sqrt{\left(r^2-(1+\rho)\right)^2 + 8r^2\rho^2}}{2}\right]^{1/2}, \tag{61}$$

$$\sigma_3(X_\rho(r)) = \left[\frac{r^2+1+\rho-\sqrt{\left(r^2-(1+\rho)\right)^2 + 8r^2\rho^2}}{2}\right]^{1/2}. \tag{62}$$

Differentiating the singular values gives

$$\frac{\partial\sigma_1}{\partial r} = 0, \tag{63}$$

$$\frac{\partial\sigma_2}{\partial r} = \frac{1}{2}\left(\frac{r^2+1+\rho+\sqrt{(r^2-(1+\rho))^2+8r^2\rho^2}}{2}\right)^{-1/2} \tag{64}$$

$$\times \frac{2r + \frac{1}{2}\left(\left((r^2-(1+\rho))^2+8r^2\rho^2\right)^{-1/2}\left(2(r^2-(1+\rho))(2r)+16r\rho^2\right)\right)}{2}, \tag{65}$$

$$\frac{\partial\sigma_3}{\partial r} = \frac{1}{2}\left(\frac{r^2+1+\rho-\sqrt{(r^2-(1+\rho))^2+8r^2\rho^2}}{2}\right)^{-1/2} \tag{66}$$

$$\times \frac{2r - \frac{1}{2}\left(\left((r^2-(1+\rho))^2+8r^2\rho^2\right)^{-1/2}\left(2(r^2-(1+\rho))(2r)+16r\rho^2\right)\right)}{2}. \tag{67}$$

Then

$$\sigma_1(X_\rho(1)) = (1-\rho)^{1/2}, \tag{68}$$

$$\sigma_2(X_\rho(1)) = (1+2\rho)^{1/2}, \tag{69}$$

$$\sigma_3(X_\rho(1)) = (1-\rho)^{1/2}, \tag{70}$$

$$\left.\frac{\partial \sigma_1}{\partial r}\right|_{r=1} = 0, \tag{71}$$

$$\left.\frac{\partial \sigma_2}{\partial r}\right|_{r=1} = \frac{1}{3}(1+2\rho)^{1/2}, \tag{72}$$

$$\left.\frac{\partial \sigma_3}{\partial r}\right|_{r=1} = \frac{2}{3}(1-\rho)^{1/2}. \tag{73}$$

Note that in fact σ_2 is the largest singular value, but for notational consistency we won't reorder.

Now we can specify g. Compute

$$\frac{\partial \ell(r)}{\partial r}|_{r=1} = g'\left((1-\rho)^{1/2}\right)\cdot 0 + g'\left((1+2\rho)^{1/2}\right)\frac{1}{3}(1+2\rho)^{1/2} + g'\left((1-\rho)^{1/2}\right)\frac{2}{3}(1-\rho)^{1/2} \tag{74}$$

$$= \frac{1}{3}g'\left((1+2\rho)^{1/2}\right)(1+2\rho)^{1/2} + \frac{2}{3}g'\left((1-\rho)^{1/2}\right)(1-\rho)^{1/2}. \tag{75}$$

Suppose $\frac{\partial \ell(r)}{\partial r}|_{r=1} = 0$. Then

$$\frac{1}{3}g'\left((1+2\rho)^{1/2}\right)(1+2\rho)^{1/2} = -\frac{2}{3}g'\left((1-\rho)^{1/2}\right)(1-\rho)^{1/2}. \tag{76}$$

What is a suitable g? We can rewrite the above constraint as

$$g'\left((1+2\rho)^{1/2}\right) = -2\left(\frac{1-\rho}{1+2\rho}\right)^{1/2} g'\left((1-\rho)^{1/2}\right). \tag{77}$$

Let $\rho = 1/2$. Then

$$g'(\sqrt{2}) = -2\left(\frac{1}{4}\right)^{1/2} g'\left((1/2)^{1/2}\right) \tag{78}$$

$$= -g'(1/\sqrt{2}) \tag{79}$$

However, by assumption, we have $g(t) = g(1/t)$, and so $g'(t) = \frac{1}{t^2}g'(1/t)$. At $t = \sqrt{2}$, $g'(\sqrt{2}) = \frac{1}{2}g'(1/\sqrt{2})$. Thus $\frac{1}{2}g'(1/\sqrt{2}) = -g'(1/\sqrt{2})$ and so $g'(1/\sqrt{2})$ and g is not monotone.

6.2.3 Proof of Proposition 9

In this proof we first show that the penalty $\|\beta\|_{1,2}$ is unchanged by unitary transformation of β.

Proposition 14 *Let $U \in \mathbb{R}^{D\times D}$ be unitary. Then $\|\beta\|_{1,2} = \|\beta U\|$.*

Proof:

$$\|\beta U\|_{1,2} = \sum_{p=1}^{P} \|\beta_{p\cdot} U\| \tag{80}$$

$$= \sum_{p=1}^{P} \|\beta_{p\cdot}\| \tag{81}$$

$$= \|\beta\|_{1,2} \tag{82}$$

□

We then show that this implies that the resultant loss is unchanged by unitary transformation of X.

Proposition 15 *Let $U \in \mathbb{R}^{D\times D}$ be unitary. Then $\widehat{\beta}(UX) = \widehat{\beta}(X)U$.*

Proof:

$$\widehat{\beta}(UX) = \arg\min_{\beta \in \mathbb{R}^{P\times D}} \|\beta\|_{1,2} \ : \ I_D = UX\beta \tag{83}$$

$$= \arg\min_{\beta \in \mathbb{R}^{P\times D}} \|\beta\|_{1,2} \ : \ U^{-1}U = U^{-1}UX\beta U \tag{84}$$

$$= \arg\min_{\beta \in \mathbb{R}^{P\times D}} \|\beta\|_{1,2} \ : \ I_D = X\beta U \tag{85}$$

$$= \arg\min_{\beta \in \mathbb{R}^{P\times D}} \|\beta U\|_{1,2} \ : \ I_D = X\beta U \tag{86}$$

$$= \arg\min_{\beta \in \mathbb{R}^{P\times D}} \|\beta\|_{1,2} \ : \ I_D = X\beta. \tag{87}$$

□

6.2.4 Proof of Proposition 10

In this proof we first show that the orthonormal submatrix $X_{.S}$ is contained within the set $\widehat{\beta}_{MBP}(w(X), I_D)$.

Proposition 16 *Let $X_{.S}$ be a rank D orthonormal submatrix of X. Then $S \in S(\beta) : \beta \in \widehat{\beta}_{MBP}(w(X), I_D))$.*

Proof: By Proposition 9 without loss of generality, let $X_{.S} = I_D$. Also, without loss of generality let $S = \{1 \dots D\}$. We show that $\beta = [I_D \ 0]$ satisfies the KKT conditions for $X = [I_D \ X_{-S}]$.

Recall that the KKT conditions for a convex optimization problem are a set of conditions on an element of the domain of the optimization algorithm that, if satisfied, certify that the element is a solution. To write out the KKT conditions, we define the Lagrangian

$$\mathcal{L}(X, \beta, \nu) = \|\beta\|_{1,2} + \nu^T(I_D - w(X)\beta) \tag{88}$$

where $\nu \in \mathbb{R}^{D\times D}$ is the dual variable.

The KKT conditions are then

- Primal feasibility: $w(X)\beta = I_D$
- Stationarity: there exists a dual variable ν such that $0 \in \partial\|\beta\|_{1,2} - w(X)^T\nu$ where ∂ is the subdifferential operator.

where dual feasibility and complementary slackness are ignorable by virtue of the absence of inequality constraints.

First note that in our case, primal feasibility is satisfied by X being rank D and w not effecting the rank of the transformed matrix since it only rescales length and not direction.

For stationarity, recall that

$$\|\beta\|_{1,2} = \sum_{p=1}^{P} \|\beta_{p.}\|_2. \tag{89}$$

Then, by the definition of subdifferential

$$\partial\|\beta\|_{1,2} = \begin{bmatrix} I_D \\ v_{d+1} \\ \dots \\ v_P \end{bmatrix} \tag{90}$$

where $v_p \in \mathbb{R}^D$ can be any vector satisfying $\|v_p\|_2 \leq 1$.

We therefore must show that there exists a ν that satisfies

$$\begin{bmatrix} I_D \\ w(X_{-S})^T \end{bmatrix} \nu = \begin{bmatrix} I_D \\ V_{-S} \end{bmatrix}. \tag{91}$$

At this point we can see that in fact I_D is an appropriate choice of ν since normalization by w leads to vectors satisfying $\|w(X)\|_2 \leq 1$.

Since for this β, $S(\beta) = S$, we have shown that S is the support of a solution to $\beta_{MBP}(w(X), I_D)$. □

The next part of the proof handles the presence of non-unique solutions.

Proposition 17 $S(\widehat{\beta}(X)) = \cup\, S(\beta) \;:\; \beta \in \beta_{MBP}(w(X), I_D)$

Proof: The proposition follows from geometric properties of the solution set. We ignore the case for singleton solutions, for which the proposition is self-evident. Recall that $\widehat{\beta}(X) := \widehat{\beta}^{\ell_2}_{MBP}(w(X), I_D) = \arg\min_{\beta \in \mathbb{R}^{P\times D}} \|\beta\|_F : \beta \in \widehat{\beta}_{MBP}(X, I_D)$. By definition of $\widehat{\beta}_{MBP}(X, I_D)$, $\|\beta_{p.}\|_2 : \beta \in \widehat{\beta}_{MBP}(X, I_D)$ is an affine linear subspace of $\mathbb{R}^P$. The projection of the vector of minimum ℓ_2 distance from the origin to a affine space intersects that space perpendicularly. However, for all $\|\beta_{.p}\|_2$ such that β is on the intersection of this affine space with the coordinate hyperplanes formed by fixing any choices of $\|\beta_{.p}\|_2 = 0$, this vector doesn't intersect this space perpendicularly. Thus, $\widehat{\beta}(X)$ is not on the intersection of the solution polytope $\widehat{\beta}_{MBP}(w(X), I_D)$ with the coordinate hyperplanes. It is therefore maximally non-sparse with respect to any of the minimizing solutions. □

6.2.5 Proof of Proposition 11

X_S admits a QR decomposition $X_S = QR$ where Q is orthogonal and R is upper triangular. By Proposition 9, we can assume without loss of generality that $Q = I_D$. That is, X_S is upper triangular. Thus, $\beta = X_S^{-1}$ is upper triangular with diagonal elements $\beta_{dd} = R^{-1}_{dd}$. $\|\beta_{d.}\|_2 \geq \beta_{dd}$, with equality only occurring when the off-diagonal elements are 0. Therefore, $\|\beta_{1,2}\| \geq D$ with equality occurring only when $\beta = I_D$.

6.2.6 Proof of Proposition 13

Proof: The augmented split Lagrangian minimized by ADMM is

$$\mathcal{L}_{\rho,\gamma,\nu}(X, \beta, z) = \|z\|_{1,2} + \mathrm{tr}(\nu^T(I_D - w(X)\beta)) + \rho\|I_D - w(X)\beta\|_F^2 + \gamma\|\beta - z\|_F^2 \tag{92}$$

This minimization problem is

$$\arg \min_{\beta, z \in P\times D} \mathcal{L}_{\rho,\gamma,\nu}(X, \beta, z) \tag{93}$$

For a more detailed explanation of the reasons to prefer optimizing the split augmented Lagrangian over direct optimization of Equation 28 see Boyd [81].

The three main subroutines (97 - 101) called by ADMM determine the per iteration complexity. The update rules of the ADMM algorithm used to optimize the augmented split Lagrangian are

$$\beta^{k+1} = \arg \min_{\beta \in P\times D} \mathcal{L}_{\rho,\gamma,\nu}(X, \beta, z^k) \tag{94}$$

$$z^{k+1} = \arg \min_{z \in P\times D} \mathcal{L}_{\rho,\gamma,\nu}(X, \beta^{k+1}, z) \tag{95}$$

$$\nu^{k+1} = \nu^k + \rho(I_D - w(X)\beta^{k+1}) \tag{96}$$

We will analyze the computational complexity of each of these steps by rewriting the update rules according to the specifics of the isometry pursuit problem.

The beta update Equation 97 projects the solution onto the feasible set.

$$\beta^{k+1} = \arg \min_{\beta \in P\times D} \mathrm{tr}(\nu^{k^T}(I_D - w(X)\beta)) + \rho\|I_D - w(X)\beta\|_2^2 + \gamma\|\beta - z^k\|_2^2 \tag{97}$$

$$= (\rho\, w(X)^T w(X) + \gamma I_D)^{-1}(\rho\, w(X)^T I_D + \gamma z^k + w(X)^T \nu^k) \tag{98}$$

The computational complexity of this computation is dominated by the inverse - that is, $O(D^2P)$.

Subsequently, the z update Equation 99 is soft-thresholding operation onto a contracted level set $\|\|\|_{1,2}$ ball.

$$z^{k+1} = \arg\min_{z \in P \times D} \|z\|_{1,2} + \gamma \|z - \beta^{k+1}\|_2^2 \tag{99}$$

$$= \begin{cases} (1 - \frac{1}{\gamma \|\beta_{j\cdot}^{k+1}\|_2})(\beta_{j\cdot}^{k+1}) \text{ if } \|\beta_{j\cdot}^{k+1}\|_2 > \frac{1}{\gamma} \\ 0 \text{ otherwise} \end{cases} \tag{100}$$

The computational complexity of this operation is $O(PD)$.

Finally, the ν update Equation 101 increases the squeezing of the z and β variables towards to each other.

$$\nu^{k+1} = \nu^k + \rho(I_D - w(X)\beta^{k+1}) \tag{101}$$

In this step, the dual variable ν that controls the penalty accrued by the splitting is updated in the direction of the constraint violation. The computational complexity of this operation is dominated by that of the matrix multiplication $w(X)\beta^{k+1}$. That is, it is $O(D^2P)$. □

6.3 Partial derivatives of candidate losses

We summarize the column-wise derivatives of the losses in Table 2. Throughout this section let $X = [x_1, \ldots, x_D] \in \mathbb{R}^{D \times D}$ be full rank. For a scalar loss $L(X)$, write

$$\nabla_X L = \left[\frac{\partial L}{\partial x_1}, \ldots, \frac{\partial L}{\partial x_D} \right]. \tag{102}$$

All formulas are stated on smooth strata: rank deficient matrices, repeated singular values, and zero rows of pseudoinverse factors should be handled using subgradients.

We will use two common identities. First, if $X = U \operatorname{diag}(\sigma) V^\top$ is a singular value decomposition and $L(X) = \sum_{d=1}^{D} \phi(\sigma_d(X))$, then

$$\nabla_X L = U \operatorname{diag}(\phi'(\sigma_1), \ldots, \phi'(\sigma_D)) V^\top. \tag{103}$$

Second, for normalized losses let $N = n_q^{\times D}(X)$. If $M = \nabla_N L$, then the column-wise chain rule is

$$\frac{\partial L}{\partial x_p} = \mathrm{D} n_q(x_p)^\top M_{\cdot p}. \tag{104}$$

For $r = \|x\|_2$, write $n_q(x) = s(r)x$ with $s(r) = q(r)/r$. Then

$$\mathrm{D} n_q(x) = s(r) I_D + \frac{s'(r)}{r} x x^\top, \tag{105}$$

$$s'(r) = \frac{q'(r)}{r} - \frac{q(r)}{r^2}, \tag{106}$$

$$q'(r) = -\frac{2e(e^r - r^{-2} e^{1/r})}{(e^r + e^{1/r})^2}. \tag{107}$$

Since $\mathrm{D} n_q(x)$ is symmetric, the transpose in this chain rule may be omitted.

Distortion. Let $\sigma_1 = \sigma_{\max}(X)$ and $\sigma_D = \sigma_{\min}(X)$ with corresponding singular vectors u_1, v_1 and u_D, v_D. For

$$L_{\mathrm{dist}}(X) = \frac{\sigma_1}{\sigma_D} - 1, \tag{108}$$

The spectral gradient identity gives

$$\nabla_X L_{\mathrm{dist}} = \frac{u_1 v_1^\top}{\sigma_D} - \frac{\sigma_1}{\sigma_D^2} u_D v_D^\top. \tag{109}$$

The derivative is insensitive to uniform rescaling, which is the sense in which this loss does not prefer unit length.

Isometry loss. For

$$L_{\text{iso}}(X) = \sum_{d=1}^{D} g(\sigma_d(X)) - D, \tag{110}$$

we have

$$g'(t) = \frac{e^t - t^{-2}e^{1/t}}{2e}. \tag{111}$$

Thus

$$\nabla_X L_{\text{iso}} = U \operatorname{diag}(g'(\sigma_1), \ldots, g'(\sigma_D))V^\top. \tag{112}$$

S-divergence to identity. Let $G = X^\top X$ and

$$L_{\text{sdiv}}(X) = \log\det\left(\frac{G + I_D}{2}\right) - \frac{1}{2}\log\det(G). \tag{113}$$

The constant $\log\det(1/2)$ does not contribute to the derivative. Using $d\log\det A = \operatorname{tr}(A^{-1}dA)$ and $dG = dX^\top X + X^\top dX$,

$$\nabla_X L_{\text{sdiv}} = 2X(G + I_D)^{-1} - XG^{-1}. \tag{114}$$

Normalized log-det volume. Let $N = n_q^{\times D}(X)$ and

$$L_{\text{nld}}(X) = -\sum_{d=1}^{D} \log \sigma_d(N). \tag{115}$$

For square full-rank N, this is $-\log|\det N|$, so

$$\nabla_N L_{\text{nld}} = -N^{-\top}. \tag{116}$$

Therefore, by the column-wise chain rule,

$$\frac{\partial L_{\text{nld}}}{\partial x_p} = -\mathrm{D}n_q(x_p)^\top (N^{-\top})_{\cdot p}. \tag{117}$$

Frobenius pseudoinverse. Let $N = n_q^{\times D}(X)$ and $B = N^{-1}$. For

$$L_{\text{frob}}(X) = \|N^{-1}\|_F - \sqrt{D}, \tag{118}$$

the differential identity $dB = -B(dN)B$ gives

$$\nabla_N L_{\text{frob}} = -\frac{N^{-\top}N^{-1}N^{-\top}}{\|N^{-1}\|_F}. \tag{119}$$

Thus

$$\frac{\partial L_{\text{frob}}}{\partial x_p} = \mathrm{D}n_q(x_p)^\top \left(-\frac{N^{-\top}N^{-1}N^{-\top}}{\|N^{-1}\|_F}\right)_{\cdot p}. \tag{120}$$

Isometry pursuit proxy. Let $N = n_q^{\times D}(X)$ and $B = N^{-1}$. The differentiable proxy for the IsometryPursuit table row is

$$L_{\text{ip}}(X) = \|N^{-1}\|_{1,2} - D. \tag{121}$$

Define $R \in \mathbb{R}^{D\times D}$ by

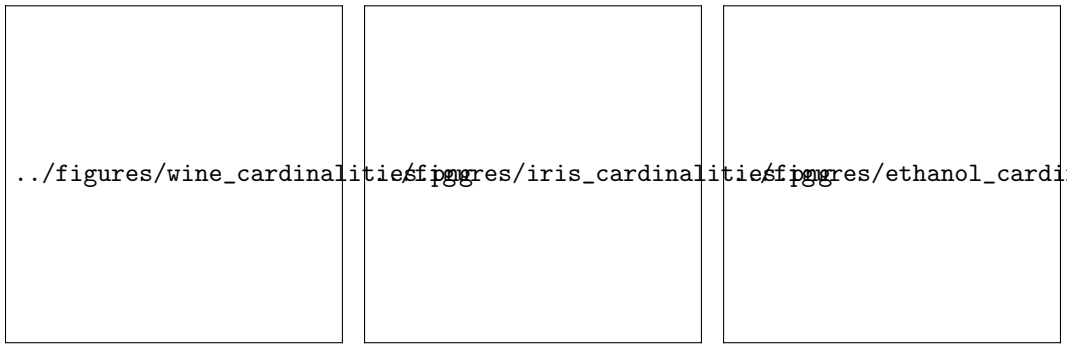
$$R_{p\cdot} = \frac{B_{p\cdot}}{\|B_{p\cdot}\|_2}. \tag{122}$$

Then $R \in \partial\|B\|_{1,2}$ on the smooth stratum where every row of B is nonzero, and

$$\nabla_N L_{\text{ip}} = -B^\top R B^\top. \tag{123}$$

The column-wise gradient is therefore

$$\frac{\partial L_{\text{ip}}}{\partial x_p} = \mathrm{D}n_q(x_p)^\top (-B^\top R B^\top)_{\cdot p}. \tag{124}$$



(a) Wine (b) Iris (c) Ethanol

Figure 6: Support cardinalities for each dataset after IsometryPursuit.

6.4 Support cardinalities

Figure 6 plots the distribution of $|\widehat{S}_{IP}|$ from Table 3 in order to contextualize the reported means. While typically $|\widehat{S}_{IP}| << P$, there are cases for Ethanol where this is not the case that drive up the means. Further investigation into the extreme failure of isometry pursuit to sparsify in the the minor chunk of the bimodal Ethanol support cardinality distribution is warranted.

6.5 Timing

While wall-time of algorithms is a non-theoretical quantity that depends on implementation details, it provides valuable context for practitioners. We therefore report the following runtimes on a 2021 Macbook Pro. The particularly high variance for brute force search in the second step of TwoStageIsometryPursuit is likely due to the large cardinalities reported in Figure 6.

Name	IP	2nd stage	Greedy
Iris	1.24 ± 0.02	0.00 ± 0.00	0.02 ± 0.00
Wine	2.32 ± 0.17	0.13 ± 0.12	0.03 ± 0.00
Ethanol	8.38 ± 0.57	0.55 ± 1.08	0.07 ± 0.01

Table 4: Algorithm runtimes in seconds across replicates.

Example 2

$$X_S = \begin{bmatrix} \frac{1}{2} & \frac{-1}{2} & \frac{1}{2} & \frac{-1}{2} & \frac{1}{2} & \frac{-1}{2} \\ \frac{\sqrt{3}}{2} & \frac{\sqrt{3}}{2} & 0 & 0 & & \\ 0 & 0 & \frac{\sqrt{3}}{2} & \frac{\sqrt{3}}{2} & 0 & 0 \\ 0 & 00 & 0 & \frac{\sqrt{3}}{2} & \frac{\sqrt{3}}{2} & \end{bmatrix}. \tag{125}$$

Proof: Let

By Given the symmetry of the second and third columns around the y axis, the loss will not change regardless of whether we upweight one or downweight the other, these changes necessarily being in conjunction. We therefore assume that their loadings with respect to the second column of the target matrix (the D dimensional identity) are the same. Clearly, they will both then have value of $\frac{1}{\sqrt{3}}$. The question then becomes, what should the values of the first column in the loading matrix β be?

We can answer this question with a simple differential equation. Let b denote the constraint violation $1 - X_{1.}\beta_{.1}$.

Then

$$\frac{\partial b}{\partial \beta_{11}} = \frac{\partial b}{\partial \beta_{21}} \tag{126}$$

$$= \frac{1}{2} \tag{127}$$

On the other hand, due to the value of $\frac{1}{\sqrt{3}}$ already assumed for the second loading of the second and third features,

$$\frac{\partial l}{\partial \beta_{11}} = \frac{\partial l}{\partial \beta_{21}} \tag{128}$$

$$= \frac{\partial}{\partial x}\sqrt{\frac{1}{3} + x^2} \tag{129}$$

$$= \frac{1}{2}(\frac{1}{3} + x^2)^{-\frac{1}{2}}(2x) \tag{130}$$

$$= \frac{x}{\sqrt{\frac{1}{3} + x^2}}. \tag{131}$$

This means that the tradeoff

$$\frac{\partial l}{\partial b} = \frac{1}{2x}\sqrt{\frac{1}{3} + x^2}. \tag{132}$$

In contrast, $\frac{\partial l}{\partial \beta_{11}} = 1$ and similarly for its constraint violation. Solving for the equilibrium, we get

$$\frac{1}{2x}\sqrt{\frac{1}{3} + x^2} = 1 \tag{133}$$

$$\frac{1}{3} + x^2 = 4x^2 \tag{134}$$

$$\frac{1}{3} = 3x^2 \tag{135}$$

$$x = \frac{1}{\sqrt{3}}. \tag{136}$$

This leads to a value for l of

$$\sqrt{\frac{1}{3} + \frac{1}{9}} + \sqrt{\frac{1}{3} + \frac{1}{9}} + \frac{2}{3} \tag{137}$$

$$= \sqrt{\frac{4}{9}} + \sqrt{\frac{4}{9}} + \frac{2}{3} \tag{138}$$

$$= \frac{2}{3} + \frac{2}{3} + \frac{2}{3} \tag{139}$$

□