Isometry pursuit

Samson Koelle
Amazon
koelle@amazon.com

Marina Meila
Department of Statistics
University of Washington
mmp@uw.edu

Abstract

Isometry pursuit is a convex algorithm for identifying orthonormal column-submatrices of wide matrices. It consists of a normalization that shrinks lengths of columns proportionally to the absolute value of the log of their length, followed by multitask basis pursuit optimization to select the longest and most orthogonal vectors. We provide theoretical and experimental results justifying this method, including a proof with realistic assumptions that such isometric submatrices, should they exist, are contained within the obtained support. For problems involving coordinate selection and diversification, Isometry Pursuit is therefore a synergistic alternative to greedy and brute force search.

1 Introduction

Many real-world problems may be abstracted as selecting a subset of the columns of a matrix representing stochastic observations or analytically exact data. This paper focuses on a simple such problem. Given a rank D matrix $X \in \mathbb{R}^{D \times P}$ with $P > D$, select a square submatrix $X_{.S}$ where subset $S \subset [P]$ satisfies $|S| = D$ that is as orthonormal as possible.

This best subset selection problem arises in applications like interpretable learning and diversification. In the former case, selection of candidate parameterizing features from an interpretable dictionary can take the above form [**????**]. While implementations vary across data and algorithmic domains, identification of such coordinates generally aids mechanistic understanding, generative control, and statistical efficiency. Diversification, rather than being concerned with selecting a uncorrelated set of features, seeks to select a dissimilar set of data points, for example in order to create balanced set of search results. The variety of measures used to assess diversity of each subset exceeds even the not exactly well-defined idea of orthonormalness. However, across cases, due to the intractibility of brute force calculations, greedy selection approaches have predominated.

Convex relaxations of best subset selection problems can provide a tractable and performant alternative to greedy optimization. Such a reformulation is said to be exact if the solution to the convex problem provides a solution to the non-convex problem. Often, a reformulation will be exact over some subset of the algorithmic input space. If an optimization problem is NP-hard, then an exact relaxation over the entire input space does not exist. Nevertheless, whether by tightness of the relaxation or likelihood of membership within this favored set, convex relaxations have enjoyed broad appeal.

This paper shows that an adapted version of the algorithm in **?** leads to a convex procedure for finding isometries. The insight leading to isometry pursuit is that multitask basis pursuit applied to an appropriately normalized X selects orthonormal submatrices. This normalization on vectors in $\mathbb{R}^D$ log-symmetrizes length and favors those closer to unit length. Multitask basis pursuit then selects a sparse orthogonal subset that are orthogonal. We provide theoretical results that formalize

[1] Work conducted outside of Amazon.

39th Conference on Neural Information Processing Systems (NeurIPS 2025).

this intuition, and give that exhibit the the usefulness of isometry pursuit as a trimming procedure prior to brute force search for diversification and interpretable coordinate selection.

2 Background

In this section, we give background on the isometry and diversification problems, as well as the spectral and convex analysis that motivate our method.

2.1 Problem

Our goal is, given a matrix $X \in \mathbb{R}^{D \times P}$, to select a subset $S \subset [P]$ with $|S| = D$ such that $X_{.S}$ is as orthonormal as possible in a computationally efficient way.

2.2 Interpretability and isometry

One motivating example is the selection of data representations from within sets of putative coordinates: the columns of a provided wide matrix. This method applies to interpretability, for which parsimony is at a premium. Interpretability arises through comparison of data with what is known to be important in the domain of the problem. This knowledge often takes the form of a functional dictionary. Evaluation of independence of dictionary features arises in numerous scenarios [**???**]. The requirement that dictionary features be full rank has been called functional independence [**?**] or feature decomposability [**?**], with connection between dictionary rank and independence via the implicit function theorem. Besides independence, the metric properties of such dictionary elements are of natural interest. This is formalized through the notion of differential.

Definition 1 *The **differential** of a smooth map $\phi : \mathcal{M} \to \mathcal{N}$ between D dimensional manifolds $\mathcal{M} \subseteq \mathbb{R}^B$ and $\mathcal{N} \subseteq \mathbb{R}^P$ is a map in tangent bases $x_1 \dots x_D$ of $T_\xi \mathcal{M}$ and $y_1 \dots y_D$ of $T_{\phi(\xi)}\mathcal{N}$ consisting of entries*

$$D\phi(\xi) = \begin{bmatrix} \frac{\partial \phi_1}{\partial x_1}(\xi) & \dots & \frac{\partial \phi_1}{\partial x_D}(\xi) \\ \vdots & & \vdots \\ \frac{\partial \phi_D}{\partial x_1}(\xi) & \dots & \frac{\partial \phi_D}{\partial x_D}(\xi) \end{bmatrix}. \tag{1}$$

It is not always necessary to explicitly estimate tangent spaces when applying this definition. The most commonly encountered manifolds are vector spaces for which the tangent spaces are trivial. This is the case for full-rank tabular data, for which isometry has a natural interpretation as a type of diversification, and often for the latent spaces of deep learning models. In this case, $B = D$.

Definition 2 *A map ϕ between D dimensional submanifolds with inherited Euclidean metric $\mathcal{M} \subseteq R^B$ and $\mathcal{N} \subseteq R^P$ ϕ is an **isometry at a point** $\xi \in \mathcal{M}$ if*

$$D\phi(\xi)^T D\phi(\xi) = I_D. \tag{2}$$

That is, ϕ is an isometry at ξ if $D\phi(\xi)$ is orthonormal.

The applications of pointwise isometry are themselves manifold. Pointwise isometric embeddings faithfully preserve high-dimensional geometry. For example, Local Tangent Space Alignment [**?**], Multidimensional Scaling [**?**] and Isomap [**?**] non-parametrically estimate embeddings that are as isometric as possible. Another approach stitches together pointwise isometries selected from a dictionary to form global embeddings [**?**]. The method is particularly relevant since it constructs such isometries through greedy search, with putative dictionary features added one at a time.

That $D\phi$ is orthonormal has several equivalent formulations. The one motivating our ground truth loss function comes from spectral analysis.

Proposition 1 *The singular values $\sigma_1 \dots \sigma_D$ are equal to 1 if and only if $U \in \mathbb{R}^{D \times D}$ is orthonormal.*

[2] Code is available at `https://github.com/sjkoelle/isometry-pursuit`.

On the other hand, the formulation that motivates our convex approach is that orthonormal matrices consist of D coordinate features whose gradients are orthogonal and of unit length.

Proposition 2 *The component vectors $u_1 \dots u_D \in \mathbb{R}^B$ form a orthonormal matrix if and only if, for all $d_1, d_2 \in [D]$, $\langle u_{d_1}, u_{d_2} \rangle = \begin{cases} 1 & d_1 = d_2 \\ 0 & d_1 \neq d_2 \end{cases}$.*

2.3 Diversification

Cosine MMR definition. Embeddings - sim is just projection

2.4 Best subset selection

Given a matrix $X \in \mathbb{R}^{D \times P}$, we compare algorithmic paradigms for solving problems of the form

$$\arg \min_{S \in \binom{[P]}{D}} l(X_{.S}) \tag{3}$$

where $\binom{[P]}{D} = \{A \subseteq [P] : |A| = D\}$. Brute force algorithms consider all possible solutions. These algorithms are conceptually simple, but have the often prohibitive time complexity $O(C_l P^D)$ where C_l is the cost of evaluating l. Greedy algorithms consist of iteratively adding one element at a time to S. This algorithms have time complexity $O(C_l PD)$ and so are computationally more efficient than brute force algorithms, but can get stuck in local minima. Formal definitions are given in Section 6.1.

Sometimes, it is possible to introduce an objective which convexifies problems of the above form. Solutions

$$\arg \min f(\beta) : Y = X\beta \tag{4}$$

to the overcomplete regression problem $Y = X\beta$ are a classic example [**?**]. When $f(\beta) = \|\beta\|_0$, this problem is non-convex, and is thus suitable for greedy or brute algorithms, but when $f(\beta) = \|\beta\|_1$, the problem is convex, and may be solved efficiently via interior-point methods. When the equality constraint is relaxed, Lagrangian duality may be used to reformulate as a so-called Lasso problem, which leads to an even richer set of optimization algorithms.

The form of basis pursuit that we apply is inspired by the group basis pursuit approach in **?**. In group basis pursuit (which we call multitask basis pursuit when grouping is dependent only on the structure of matrix-valued response variable y) the objective function is $f(\beta) = \|\beta\|_{1,2} := \sum_{p=1}^{P} \|\beta_{p.}\|_2$ [**???**]. This objective creates joint sparsity across entire rows $\beta_{p.}$ and was used in **?** to select between sets of interpretable features.

3 Method

We adapt the group lasso paradigm used to select independent dictionary elements in **??** to select pointwise isometries from a dictionary. We first define a ground truth objective computable via brute and greedy algorithms that is uniquely minimized by orthonormal matrices. We then define the combination of normalization and multitask basis pursuit that approximates this ground truth loss function. We finally give a brute post-processing method for ensuring that the solution is D sparse, and provide a theoretical result that the two stage approach will always result in recovery of an isometric solution from the dictionary, should one exist.

3.1 Ground truth

We first introduce a novel ground truth objective function against which we measure the success of our algorithm. We'd like a ground truth objective that is minimized uniquely by orthonormal matrices, invariant under rotation, and depends on all the singular values of the matrix. Deformation [**?**] uses only the extremal eigenvalues, while nuclear norm and log-determinant [**?**] and are not uniquely minimized at unitarity. We therefore introduce an alternative ground truth objective that satisfies the above desiderata. We will use this objective as the ground truth against which we evaluate our approach.

This ground truth objective is

$$l_c : \mathbb{R}^{D\times P} \to \mathbb{R}^+ \tag{5}$$

$$X \mapsto \sum_{d=1}^{D} g(\sigma_d(X), c) \tag{6}$$

where $\sigma_d(X)$ is the d-th singular value of X and

$$g : \mathbb{R}^+ \times \mathbb{R}^+ \to \mathbb{R}^+ \tag{7}$$

$$t, c \mapsto \frac{e^{t^c} + e^{t^{-c}}}{2e}. \tag{8}$$

By Proposition 1, we can see that l_c is uniquely maximized by orthonormal matrices. Moreover, g is convex, and $l_c(X^{-1}) = l_c(X)$ when X is invertible. Figure 1 gives a graph of l_c when $D = 1$.

Our ground truth objective is therefore the best subset selection problem

$$\arg\min_{S\in\binom{[P]}{d}} l_c(X_{.S}). \tag{9}$$

Regardless of the convexity of l_c, brute combinatorial search over $[P]$ is inherently non-convex. This motivates our alternative formulation.

3.2 Normalization

Since basis pursuit methods tend to select longer vectors, selection of orthonormal submatrices requires normalization such that both long and short candidate basis vectors are penalized in the subsequent regression. We introduce the following definition.

Definition 3 (Symmetric normalization) *A function $q : \mathbb{R}^D \to \mathbb{R}^+$ is a symmetric normalization if*

$$\arg\max_{v\in\mathbb{R}^D} q(v) = \{v : \|v\|_2 = 1\} \tag{10}$$

$$q(v) = q(\frac{v}{\|v\|_2^2}) \tag{11}$$

$$q(v_1) = q(v_2)\ \forall\ v_1, v_2 \in \mathbb{R}^D : \|v_1\|_2 = \|v_2\|_2. \tag{12}$$

We use such functions to normalize column-vector length in such a way that vectors of length 1 prior to normalization have longest length after normalization and vectors are shrunk proportionately to their deviation from 1. That is, we normalize vectors by

$$n : \mathbb{R}^D \to \mathbb{R}^D \tag{13}$$

$$v \mapsto q(v)v \tag{14}$$

and matrices by

$$w : \mathbb{R}^{D\times P} \to \mathbb{R}^D \tag{15}$$

$$X_{.p} \mapsto n(X_{.p})\ \forall\ p \in [P]. \tag{16}$$

Given $c > 0$, we choose q as follows.

$$q_c : \mathbb{R}^D \to \mathbb{R}^+ \tag{17}$$

$$v \mapsto \frac{e^{\|v\|_2^c} + e^{\|v\|_2^{-c}}}{2e}. \tag{18}$$

Besides satisfying the conditions in Definition 3, this normalization has some additional nice properties. First, q is convex. Second, it grows asymptotically log-linearly. Third, while $\exp(-|\log t|) = \exp(-\max(t, 1/t))$ is a seemingly natural choice for normalization, it is non smooth, and the LogSumExp [**?**] replacement of $\max(t, 1/t)$ with $\log(\exp(t) + \exp(1/t))$ simplifies to 18 upon exponentiation. Finally, the parameter c grants control over the width of the basin, which may be useful for avoiding numerical issues arising close to 0 and ∞.

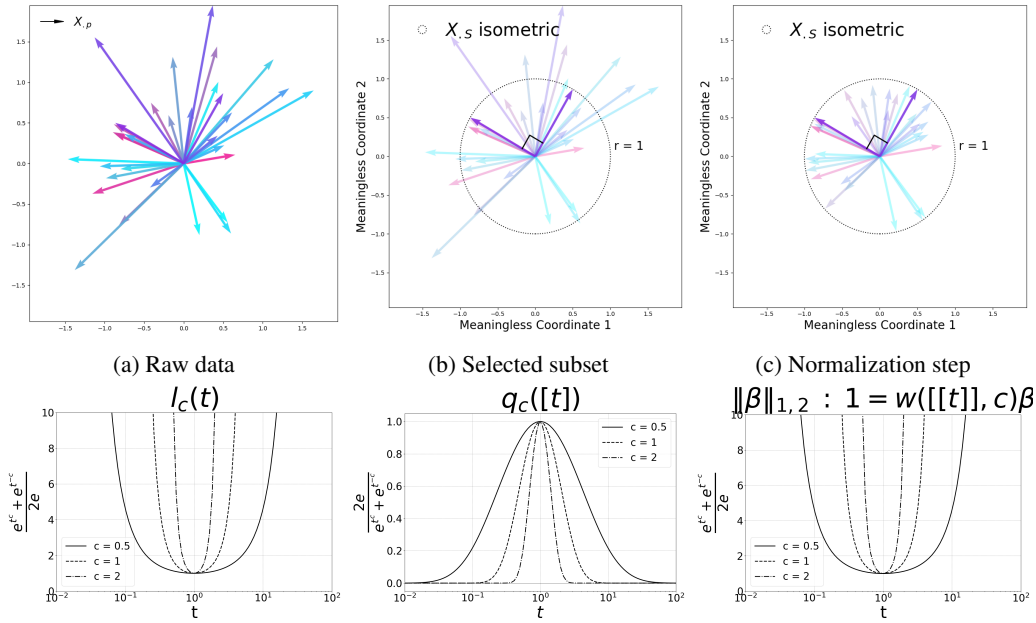

(a) Raw data (b) Selected subset (c) Normalization step

(d) Ground truth loss (e) Normalized length vs. unnormalized length (f) Basis pursuit loss

Figure 1: 1a raw data X represented as column-vectors. 1b a isometric subset - the identification of which is our objective. 1c the vectors after normalization so that vectors of length 1 maintain longest length after normalization. 1d our ground truth loss function. 1e normalized length used to rescale Figure 1a to get vectors in Figure 1c. 1fd basis pursuit loss for different values of c in the one-dimensional case $D = 1$. This loss is equivalent to that shown in Figure 1d in the one-dimensional case.

3.3 Isometry pursuit

Isometry pursuit is the application of multitask basis pursuit to the normalized design matrix $w(X, c)$ to identify submatrices of X that are as orthonormal as possible. Define the multitask basis pursuit penalty

$$\|\cdot\|_{1,2} : \mathbb{R}^{P\times D} \to \mathbb{R}^{+} \tag{19}$$

$$\beta \mapsto \sum_{p=1}^{P} \|\beta_{p.}\|_2. \tag{20}$$

Given a matrix $Y \in \mathbb{R}^{D\times D}$, the multitask basis pursuit solution is

$$\widehat{\beta}_{MBP}(X, Y) := \arg \min_{\beta\in\mathbb{R}^{P\times D}} \|\beta\|_{1,2} \ : \ Y = X\beta. \tag{21}$$

Isometry pursuit is then given by

$$\widehat{\beta}_c(X) := \widehat{\beta}_{MBP}(w(X, c), I_D) \tag{22}$$

where I_D is the D dimensional identity matrix and recovered functions are the indices of the dictionary elements with non-zero coefficients. That is, they are given by $S(\beta)$ where

$$S : \mathbb{R}^{P\times D} \to \binom{[P]}{D} \tag{23}$$

$$\beta \mapsto \{p \in [P] : \|\beta_{p.}\| > 0\}. \tag{24}$$

ISOMETRYPURSUIT(Matrix $X \in \mathbb{R}^{D\times P}$, scaling constant c)

1: Normalize $X_c = w(X, c)$
2: Optimize $\widehat{\beta} = \widehat{\beta}_{MBP}(X_c, I_D)$
3: **Output** $\widehat{S} = S(\widehat{\beta})$

3.4 Theory

The intuition behind our application of multitask basis pursuit is that submatrices consisting of vectors which are closer to 1 in length and more orthogonal will have smaller loss. A key initial theoretical assertion is that ISOMETRYPURSUIT is invariant to choice of basis for X.

Proposition 3 *Let $U \in \mathbb{R}^{D \times D}$ be orthonormal. Then $S(\widehat{\beta}(UX)) = S(\widehat{\beta}(X))$.*

A proof is given in Section 6.2.1. This has as an immediate corollary that we may replace I_D in the constraint by any orthonormal $D \times D$ matrix.

When a rank D orthonormal column-submatrix X_S exists, a minimizer of Program 21 consists of the inverse of this submatrix.

Proposition 4 *Let w_c be a normalization satisfying the conditions in Definition 3. Let $X \in \mathbb{R}^{D \times P}$ contain an orthonormal $D \times D$ matrix X_S. Then $S \in \mathcal{S}_{IP}(X)$.*

However, this solution is not necessarily unique.

Proposition 5

Proofs of these propositions are given in Section 6.2.2, with corresponding experimental results in Section 4.

Non-unique solutions to 22 are subtly more common than in standard lasso theory. Recall that, given a design matrix X, there exists a response variable y such that the lasso admits a non-unique solution if and only if the rowspan of the design matrix X contains a sufficient codimension edge of the p dimensional hypercube [**?**]. This condition generalizes the previously introduced general position condition - that affinely independent columns of the design matrix guarantee non-uniqueness. However, for Isometry Pursuit, non-unique solutions may occur even when this condition is satisfied. A simple example - $X = \begin{bmatrix} 1 & 0 & \frac{\sqrt{2}}{2} & \frac{\sqrt{2}}{2} \\ 0 & 1 & \frac{\sqrt{2}}{2} & \frac{\sqrt{-2}}{2} \end{bmatrix}$ results from the rotation invariance Proposition 4, but more subtle examples exist.

3.5 Optimization

A typical response to considerations of non-uniqueness in such problems is the characterization of the "min-norm" solution - that is, the solution which minimizes the ℓ_2 norm. This implicit regularization assumption is reasonable. For example, in **?** the implicit selection of the minimum l2-norm solution among the lasso solutions was proven for the Least Angle Regression Selection (LARS) algorithm for solving the lasso, while **?** simply assumes that the group lasso solution is in fact min-norm prior to subsequent theoretical analyses. Empirically, the Splitting Conic Solver type method (which is itself a variant of the Alternating Direction Method of Multilpliers) stably estimates the min-norm solution when initialized at 0, but we leave theoretical and experimental proof of this feature of this feature our optimization approach for future work, and instead make the ability of our optimizer to obtain the min-norm solution an assumption of our final proposition.

Proposition 6 *Let $S_{IP}^{\ell_2}(X)$ be the min-norm Isometry Pursuit solution. Let w_c be a normalization satisfying the conditions in Definition 3. Let $X \in \mathbb{R}^{D \times P}$ contain an orthonormal $D \times D$ matrix X_S. Then $S \subseteq S_{IP}^{\ell_2}(X)$.*

3.6 Two-stage isometry pursuit

Proposition 6 suggests the following algorithm, which first uses the convex problem to prune and then apply brute search upon the substantially reduced feature set. Similar two-stage approaches are standard in the Lasso literature **??**. This method forms our practical isometry estimator, and is discussed further in Sections 5 and 6.4.

TWOSTAGEISOMETRYPURSUIT(Matrix $X \in \mathbb{R}^{D \times P}$, scaling constant c)

1: $\widehat{S}_{IP} = \text{ISOMETRYPURSUIT}(X, c)$
2: $\widehat{S} = \text{BRUTESEARCH}(X_{.\widehat{S}_{IP}}, l_c)$
3: **Output** $\widehat{S}$

This approach is guaranteed to select an isometric submatrix should one exist. In fact, it will identify the brute force minimizer of the corresponding lasso dual problem. Note however that if an isometric column submatrix does not exist, then it will not identify the optimum of Program 9, since the lasso objective and the singular value objective do not necessarily share a minimizer for all choices of X.

4 Experiments

Say you are hosting an elegant dinner party, and wish to select a balanced set of wines for drinking and flowers for decoration. We demonstrate TWOSTAGEISOMETRYPURSUIT and GREEDYSEARCH on the Iris and Wine datasets [**???**]. This has an intuitive interpretation as selecting diverse elements that reflects the peculiar structure of the diversification problem. Features like *petal width* are rows in X. They are features on the basis of which we may select among the flowers those which are most distinct from another. Thus, in diversification, $P = n$.

We also analyze the Ethanol dataset from **??**, but rather than selecting between bourbon and scotch we evaluate a dictionary of interpretable features - bond torsions - for their ability to parameterize the molecular configuration space. In this interpretability use case, columns denote gradients of informative features. We compute Jacoban matrices of putative parametrization functions and project them onto estimated tangent spaces (see **?** for preprocessing details). Rather than selecting between data points, we are selecting between functions which parameterize the data.

For basis pursuit, we use the SCS interior point solver [**?**] from CVXPY [**??**], which is able to push sparse values arbitrarily close to 0 [**?**]. Statistical replicas for Wine and Iris are created by resampling across $[P]$. Due to differences in scales between rows, these are first standardized. For the Wine dataset, even BRUTESEARCH on $\widehat{S}_{IP}$ is prohibitive in $D = 13$, and so we truncate our inputs to $D = 6$. For Ethanol, replicas are created by sampling from data points and their corresponding tangent spaces are estimated in $B = 252$.

Figure 2 and Table 1 show that the l_1 accrued by the subset $\widehat{S}_G$ estimated using GREEDYSEARCH with objective l_1 is higher than that for the subset estimated by TWOSTAGEISOMETRYPURSUIT. This effect is statistically significant, but varies across datapoints and datasets. Figure 3 details intermediate support recovery cardinalities from ISOMETRYPURSUIT. We also evaluated second stage BRUTESEARCH selection after random selection of $\widehat{S}_{IP}$ but do not report it since it often lead to catastrophic failure to satisfy the basis pursuit constraint. Wall-clock runtimes are given in Section 6.5.

Name	D	P	R	c	$l_1(X_{.\widehat{S}_G})$	$\lvert\widehat{S}_{IP}\rvert$	$l_1(X_{.\widehat{S}})$	$P_R(l_1(X_{.\widehat{S}_G}) > l_1(X_{.\widehat{S}}))$	$P_R(l_1(X_{.\widehat{S}_G}) = l_1(X_{.\widehat{S}}))$	$\widehat{P}(\bar{l}_1(X_{.\widehat{S}_G}) > \bar{l}_1(X_{.\widehat{S}}))$
Iris	4	75	25	1	13.8 ± 7.3	7 ± 1	6.9 ± 1.4	0.96	0.	2.4e-05
Wine	6	89	25	1	7.7 ± 0.3	13 ± 2	7.6 ± 0.3	0.64	0.16	6.3e-04
Ethanol	2	756	100	1	2.6 ± 0.3	90 ± 165	2.5 ± 0.2	0.66	0.17	2.1e-05

Table 1: Experimental parameters and results. For Iris and Wine, P results from random downsampling by a factor of 2 to create R replicates. P_R values are empirical probabilities, while estimated P-values $\widehat{P}$ are computed by paired two-sample T-test on $l_1(X_{.\widehat{S}})$ and $l_1(X_{.\widehat{S}_G})$. For brevity, in this table $\widehat{S} := \widehat{S}_{TSIP}$.

5 Discussion

We have shown that multitask basis pursuit can help select isometric submatrices from appropriately normalized wide matrices. This approach - isometry pursuit - is a convex alternative to greedy

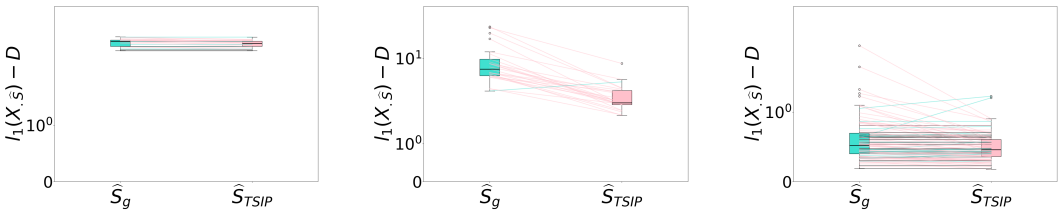

(a) Wine dataset (b) Iris dataset (c) Ethanol dataset

Figure 2: Isometry losses l_1 for Wine, Iris, and Ethanol datasets across R replicates. Lower greedy losses are shown with turquoise, while lower two stage losses are shown with pink. Equal losses are shown with black lines. As detailed in Table 1, losses are generally lower for two-stage isometry pursuit solutions.

methods for selection of orthonormalized features from within a dictionary. Isometry pursuit can be applied to diversification and geometrically-faithful coordinate estimation. Our experiments exemplify these applications, but more can be done. One potential application is diversification in recommendation systems [**???**] and other retrieval systems such as in RAG [**?????**]. Another is decomposing interpretable yet overcomplete dictionaries in transformer residual streams, with each token considered as generating its own tangent space [**??**].

Compared with the greedy algorithms used in such areas [**??????????**], the convex reformulation may add speed and convergence to a global minima. The comparison of greedy **????** and convex [**???**] basis pursuit formulations has a rich history, and theoretical understanding of the behavior of this approximation is evolving. Diversification problems have been cited as NP-hard, and isometry pursuit can be considered analogous to them in the sense of basis pursuit and the lasso against best subset selection, with the caveat that best subset selection of the basis pursuit loss minimizer isn't totally equivalent to isometry pursuit even though they share the same unique optimum. Characterization of solutions resulting from removal of the restriction $P = D$ on the conditions of Proposition 4 may help justify the second selection step. That the solution of a lasso problem can sometimes be a non-singleton set is well-known [**???????????**]. Perhaps surprisingly, it appears empirically that for isometry pursuit that this can occur even when the design matrix is not in general position.

This convex set appears to contain the brute solution of a related problem. The convergence of SCS algorithm to the 2-norm minimizing solution due to the Lagrangian dual constraint penalty and the convexity of the loss minimizer preimage suggest that a related two stage procedure always succeeds in identifying the brute $\|\|_{1,2}$ minimizer. Related conditions have been discussed in **??**, and we examine this topic experimentally in Section 6.4.

Algorithmic variants include the multitask lasso [**?**] extension of our estimator, as well as characterization of D function selection within $\mathbb{R}^B$. Tangent-space specific variants have been studied in more detail in **??** with additional grouping across datapoints, and a corresponding variant of the isometry theorem that missed non-uniqueness was claimed in **?**. Comparison of our loss with curvature - whose presence prohibits D element isometry - could prove fertile, as could comparison with the so-called restricted isometry property used to show guaranteed recovery at fast convergence rates in supervised learning [**??**].

6 Supplement

This section contains algorithms, proofs, and experiments in support of the main text.

6.1 Algorithms

We give definitions of the brute and greedy algorithms for the combinatorial problem studied in this paper. The brute force algorithm is computationally intractable for all but the smallest problems, but always finds the global minima.

BRUTESEARCH(Matrix $X \in \mathbb{R}^{D \times P}$, objective f)

for each combination $S \subseteq \{1, 2, \ldots, P\}$ with $|S| = D$ **do**
  Evaluate $f(X_{.S})$
end for
Output the combination S^* that minimizes $f(X_{.S})$

Greedy algorithms are computationally expedient but can get stuck in local optima [**??**], even with randomized restarts [**?**].

GREEDYSEARCH(Matrix $X \in \mathbb{R}^{D \times P}$, objective f, selected set $S = \emptyset$, current size $d = 0$)

if $d = D$ **then**
  Return S
else
  Initialize $S_{\text{best}} = S$
  Initialize $f_{\text{best}} = \infty$
  for each $p \in \{1, 2, \ldots, P\} \setminus S$ **do**
    Evaluate $f(X_{.(S \cup \{p\})})$
    if $f(X_{.(S \cup \{p\})}) < f_{\text{best}}$ **then**
      Update $S_{\text{best}} = S \cup \{p\}$
      Update $f_{\text{best}} = f(\mathcal{X}_{.(S \cup \{p\})})$
    end if
  end for
  Return GREEDYSEARCH(X, f, S_{best}, $d + 1$)
end if

6.2 Proofs

6.2.1 Proof of Proposition 3

In this proof we first show that the penalty $\|\beta\|_{1,2}$ is unchanged by unitary transformation of β.

Proposition 7 *Let $U \in \mathbb{R}^{D \times D}$ be unitary. Then $\|\beta\|_{1,2} = \|\beta U\|$.*

Proof:

$$\|\beta U\|_{1,2} = \sum_{p=1}^{P} \|\beta_{p.} U\| \tag{25}$$

$$= \sum_{p=1}^{P} \|\beta_{p.}\| \tag{26}$$

$$= \|\beta\|_{1,2} \tag{27}$$

□

We then show that this implies that the resultant loss is unchanged by unitary transformation of X.

Proposition 8 *Let $U \in \mathbb{R}^{D \times D}$ be unitary. Then $\widehat{\beta}(UX) = \widehat{\beta}(X)U$.*

Proof:

$$\widehat{\beta}(UX) = \arg \min_{\beta \in \mathbb{R}^{P \times D}} \|\beta\|_{1,2} \ : \ I_D = UX\beta \tag{28}$$

$$= \arg \min_{\beta \in \mathbb{R}^{P \times D}} \|\beta\|_{1,2} \ : \ U^{-1}U = U^{-1}UX\beta U \tag{29}$$

$$= \arg \min_{\beta \in \mathbb{R}^{P \times D}} \|\beta\|_{1,2} \ : \ I_D = X\beta U \tag{30}$$

$$= \arg \min_{\beta \in \mathbb{R}^{P \times D}} \|\beta U\|_{1,2} \ : \ I_D = X\beta U \tag{31}$$

$$= \arg \min_{\beta \in \mathbb{R}^{P \times D}} \|\beta\|_{1,2} \ : \ I_D = X\beta. \tag{32}$$

□

6.2.2 Proof of Proposition 4

Proposition 9 *Let w_c be a normalization satisfying the conditions in Definition Let $X \in \mathbb{R}^{D \times P}$ contain a rank D orthonormal submatrix $X_{.S} \in \mathbb{R}^{D \times D}$. Then $S \subseteq S(\{\arg\min_{\beta \in \mathbb{R}^{D \times P}} \|\beta\|_{1,2} : I_D = \ w(X,c)\beta\} \ D)$.*

Proof: Without loss of generality, we show that $[I_D 0]$ satisfies the KKT conditions for $X = [I_D X_{-S}]$ The KKT conditions are

- Primal feasibility: $w(X,c)\beta = I_D$
- Stationarity: there exists a dual variable $\nu \in \mathbb{R}^{D \times D}$ such that $0 \in \partial\|\beta\|_{1,2} - w(X,c)^T\nu$ where ∂ is the subdifferential operator.
- Dual feasibility
- Complementary slackness

Stationarity:

$$\|\beta\|_{1,2} = \sum_{p=1}^{P} \|\beta_{p.}\|_2. \tag{33}$$

Then

$$\partial\|\beta\|_{1,2} = \begin{bmatrix} I_D v_{d+1} \\ \cdots \\ v_P \end{bmatrix} \tag{34}$$

where $v_p \in \mathbb{R}^D$ satisfies $\|v_p\|_2 \leq 1$.

We therefore must satisfy

$$\begin{bmatrix} I_D \\ w(X_{-S}, c)^T \end{bmatrix} \nu = \begin{bmatrix} I_D \\ V_{-S} \end{bmatrix} \tag{35}$$

where $\nu \in \mathbb{R}^{D \times D}$.

At this point we can see that in fact I_D is an appropriate choice of ν by the boundary on the range of the normalized vector.

□

Proposition 10 *The SCS algorithm identities the minimum 2-norm solution*

Proof: The splitting conic solver algorithm, also known as alternative direction method of multipliers (ADMM), makes use of the follow "split augmented Lagrangian" objective function

$$\mathcal{L}_\rho(X, Z\beta, \lambda_1) = \|\beta\|_{1,2} + \lambda_1^T(I_D - X\beta) + \lambda_2^T(\beta - Z) + \rho\|I_D - X\beta\|_2^2 + \rho\|\beta - Z\|_2^2 \tag{36}$$

This function provably has the same minimizer over β as **??**, but is more amenable to optimization. By differentiating this expression w.r.t. β to write out its update rule, we show that the update requires computing a pseudoinverse. This pseudoinverse is in fact chosen to be a ridge regression pseudoinverse (i.e. $(X^TX + \Lambda)^{-1}$) which, as the closed form solution for ridge regression, for $\|\beta\|_2^2$ minimization.

Proposition 11 *Equation 36 is strongly convex.*

Proof: □

Since the squared 2-norm is strongly convex, this minimizer is unique. □

Proposition 12 *The minimizing solution is contained within the estimated set*

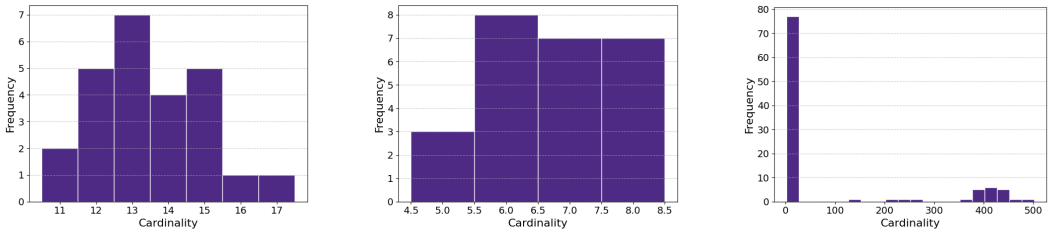

(a) Wine Dataset (b) Iris Dataset (c) Ethanol Dataset

Figure 3: Support Cardinalities for Wine, Iris, and Ethanol datasets

6.3 Support cardinalities

Figure 3 plots the distribution of $|\widehat{S}_{IP}|$ from Table 1 in order to contextualize the reported means. While typically $|\widehat{S}_{IP}| << P$, there are cases for Ethanol where this is not the case that drive up the means.

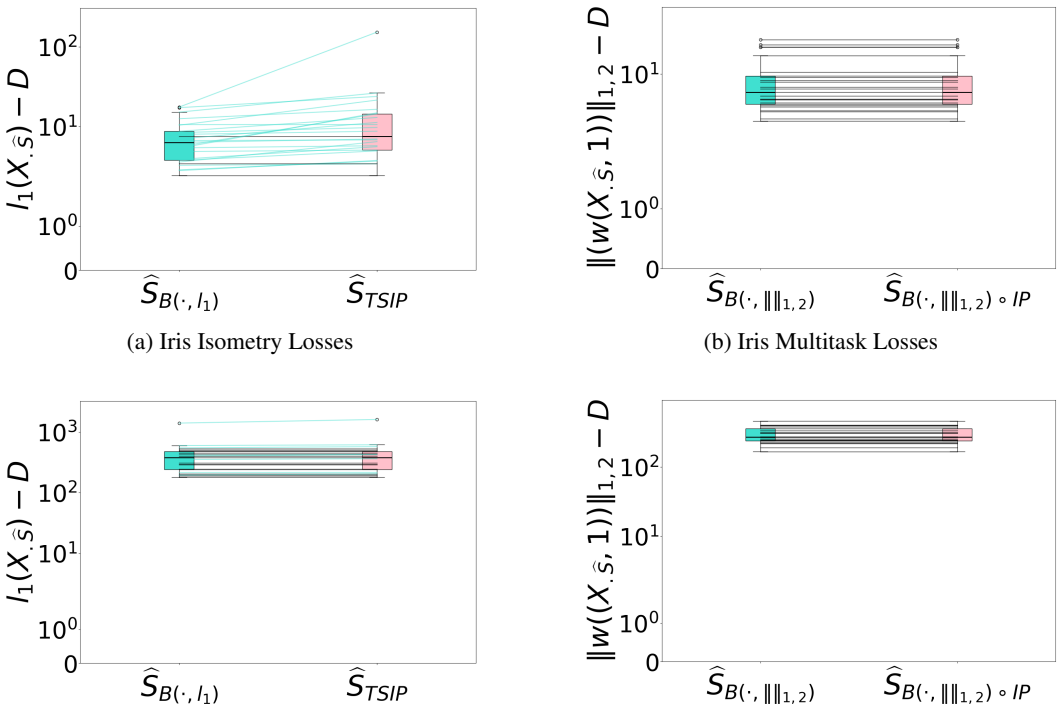

(a) Iris Isometry Losses

(b) Iris Multitask Losses

(c) Wine Isometry Losses

(d) Wine Multitask Losses

Figure 4: Comparison of Isometry and Group Lasso Losses across 25 replicates for randomly downsampled Iris and Wine Datasets with $(P, D) = (4, 15)$ and $(13, 18)$, respectively. Note that this further downsampling compared with Section 4 was necessary to compute global minimizers of BRUTESEARCH. Lower brute losses are shown with turquoise, while lower two stage losses are shown with pink. Equal losses are shown with black lines.

6.4 Proposition 4 deep dive

As mentioned in Section 5, the conditions under which the restriction $P = D$ in Proposition 4 may be relaxed are of theoretical and practical interest. The results in Section 4 show that there are circumstances in which the GREEDYSEARCH performs better than TWOSTAGEISOMETRYPURSUIT, so clearly TWOSTAGEISOMETRYPURSUIT does not always achieve a global optimum. Figure 4 gives results on the line of inquiry about why this is the case based on the reasoning presented in Section 5. In these results a two-stage algorithm achieves the global optimum of a slightly different brute problem, namely brute optimization of the multitask basis pursuit penalty $\|\cdot\|_{1,2}$. That is, brute search on $\|\cdot\|_{1,2}$ gives the same result as the two stage algorithm with brute search on $\|\cdot\|_{1,2}$ subsequent to isometry pursuit. This suggests that failure to select the global optimum by TWOSTAGEISOMETRYPURSUIT is in fact only due to the mismatch between global optimums of brute optimization of the multitask penalty and the isometry loss given certain data. Theoretical formalization, as well as investigation of what data configurations this equivalence holds for, is a logical follow-up.

6.5 Timing

While wall-time of algorithms is a non-theoretical quantity that depends on implementation details, it provides valuable context for practitioners. We therefore report the following runtimes on a 2021 Macbook Pro. The particularly high variance for brute force search in the second step of TWOSTAGEISOMETRYPURSUIT is likely due to the large cardinalities reported in Figure 3.

Name	IP	2nd stage brute	Greedy
Iris	1.24 ± 0.02	0.00 ± 0.00	0.02 ± 0.00
Wine	2.32 ± 0.17	0.13 ± 0.12	0.03 ± 0.00
Ethanol	8.38 ± 0.57	0.55 ± 1.08	0.07 ± 0.01

Table 2: Algorithm runtimes in seconds across replicates.